

PIPE-Cypher: Automatic Enterprise Benchmark Generation for Text-to-Cypher Systems

Anonymous ACL submission

Abstract

Enterprises increasingly use property graphs and Cypher to analyze fraud, risk, identity, customer, and operational data. A useful enterprise Text2Cypher benchmark must therefore be private, refreshable, executable, and sensitive to graph-specific semantics such as relationship direction, literal grounding, and schema vocabulary. We present PIPE-Cypher, a local-model pipeline for generating organization-specific natural-language-to-Cypher benchmarks. PIPE-Cypher combines schema introspection, reverse-query grounding, constrained generation, deterministic Cypher governance, execution validation, diversity controls, and LLM-judge review. In live runs with local Qwen3.5-9B, PIPE-Cypher exports 3,000 accepted FinBench/SNB examples, completes three audited ablation suites, calibrates an automated judge against human labels, and onboards ICIJ Offshore Leaks as a third public finance/compliance graph.

1 Introduction

Property graphs are attractive in enterprise settings because the facts of interest are often relational paths: account transfers, identity entitlements, access chains, customer interactions, supplier dependencies, and fraud rings. Cypher gives analysts a compact language for these patterns. As LLMs become natural-language interfaces for graph analytics, an organization cannot evaluate a Text2Cypher system only on public schemas; it needs to know whether the model handles its labels, relationship directions, values, governance rules, and recurring operational questions.

This changes the benchmark problem. A static public dataset is valuable for shared comparison, but it cannot contain a bank’s account taxonomy, an identity team’s permission graph, or the exact categorical values that make a query answerable. It also cannot refresh when schemas change. Industry teams therefore need a benchmark factory: a

repeatable way to turn a live property graph into a balanced, executable, privacy-aware NL-to-Cypher benchmark without sending sensitive schema or values to paid generation APIs.

PIPE-Cypher addresses this setting. The pipeline profiles a target graph, grounds question slots through reverse Cypher execution, constrains local-model generation, validates and repairs generated Cypher through deterministic gates, and uses a local LLM judge only after execution evidence is available. Its design treats Cypher correctness as a governed artifact rather than a prompt preference: relationship direction, read-only safety, exact literal use, categorical values, contextual return columns, and `RETURN DISTINCT` are checked or normalized before examples are accepted.

Contributions. We make four contributions: (1) an end-to-end local-model pipeline for private enterprise NL-to-Cypher benchmark generation; (2) a Cypher governance layer that turns production query-writing constraints into deterministic validation and conservative rewrite rules; (3) a scaled evaluation over FinBench, SNB, and ICIJ showing balanced generation, ablations, judge calibration, and downstream model utility; and (4) reproducible artifacts for enterprise onboarding, privacy redaction, value sampling, benchmark refresh, and appendix-level auditability.

2 Related Work

Recent Text2Cypher resources, including Mind the Query (Chauhan et al., 2025), SyntheT2C (Zhong et al., 2025), Auto-Cypher (Tiwari et al., 2025), and Text2Cypher (Ozsoy et al., 2025), show that high-quality Cypher data requires schema grounding, execution checks, verification, and complexity-aware evaluation. Mind the Query is especially relevant because it reports an Industry Track Text2Cypher dataset with schema, runtime, value, and human logical validation. PIPE-Cypher borrows the

reviewer-facing discipline of those checks, but changes the target artifact: instead of releasing one static dataset, it builds a rerunnable generator for an organization’s own graph, local model endpoint, privacy policy, and benchmark refresh cycle.

Text-to-SQL benchmarks such as Spider 2.0 (Lei et al., 2024) and BIRD (Li et al., 2023) have pushed text-to-query evaluation toward realistic database tasks and execution-based scoring. Graph query generation adds different failure modes: relationship direction can invert the meaning of a query, node and relationship properties live in different namespaces, and path-shaped questions can be syntactically valid while semantically ungrounded. CIKM AutoQuery (Zheng et al., 2024) motivates treating workload generation as a separate object of study from downstream model quality. LDBC FinBench (Qi et al., 2023) and SNB (Puroja et al., 2023) provide public graph workloads with financial and social-network structure, while ICIJ Offshore Leaks gives an additional public finance/compliance onboarding check.

For benchmark quality, lexical diversity alone is not enough. PIPE-Cypher combines text-generation diagnostics such as Distinct-n (Li et al., 2016) and self-BLEU-style redundancy checks (Zhu et al., 2018) with text-to-query structure metrics: schema coverage, relationship/property coverage, query-signature diversity, structural feature rates, and normalized entropy over graph/category/difficulty cells. These metrics make a practical distinction that matters to enterprise users: a benchmark can be syntactically diverse while still overusing the same values or query templates.

3 Method

PIPE-Cypher has six stages: schema profiling, workload planning, reverse Cypher grounding, constrained Cypher generation and repair, deterministic validation and execution, and LLM-judge review. The central design choice is to make answerability and governance executable. Prompts can ask a model to respect relationship directions or exact literals, but accepted examples must prove those properties through schema checks, parser-style structure extraction, live execution, and judge review.

Schema profiling records labels, relationship types, properties, observed directions, and bounded low-cardinality categorical values. Workload plan-

ning targets eight operational categories: simple retrieval, complex retrieval, simple aggregation, complex aggregation, boolean existence, negation/difference, path/temporal transaction, and ranking/top-k. Reverse grounding runs read-only Cypher to bind slots to values that actually produce rows, reducing the common synthetic-data failure where a plausible question has no answer in the graph.

The deterministic layer enforces read-only safety, syntax shape, schema-visible labels and properties, explicit relationship types, observed relationship directions, schema-provided categorical values, and non-empty execution where required. A lightweight Cypher analyzer extracts return aliases, variables, labels, relationship observations, risky constructs, and rewrite skip reasons so normalization is auditable rather than a silent string edit. The direction gate interprets both outgoing and incoming Cypher arrow syntax before schema checking, and rejects undirected relationship patterns so directionality is an executable constraint rather than only a prompt instruction.

4 Implementation

The implementation is a Python package with Neo4j as the experimental backend. The paper frames the contribution around Cypher and property graphs rather than a single vendor. All reported generation, judging, and downstream evaluation runs use a local Qwen3.5-9B endpoint behind a vLLM/OpenAI-compatible interface. This keeps the study aligned with an enterprise deployment pattern in which the benchmark pipeline runs inside the organization’s compute boundary without paid generation APIs.

The built-in FinBench profile is grounded in the public datagen snapshot export and includes typed node and relationship properties. Live schema inspection also records low-cardinality string values as categorical constraints for prompting and validation. The generated Neo4j import script uses uniqueness constraints for nodes and creates, rather than merges, transaction relationships so repeated account-to-account events are preserved. The SNB reference profile is grounded in the official Neo4j/Cypher headers and read-query files.

For generation, PIPE-Cypher uses seeded workload templates with deterministic reverse-binding queries and a mixed mode that prepends proven workload seeds to LLM-proposed templates. Ev-




Figure 1: PIPE-Cypher benchmark generation pipeline. The key industry additions beyond static Text2Cypher dataset construction are privacy/value policies, reverse grounding, Cypher governance, execution diagnostics, local judge calibration, and benchmark export/refresh.

Gate	Evidence checked	Failure mode caught
Read-only safety	blocked Cypher write/admin tokens	destructive or operational queries
Schema validity	labels, relationship types, properties, categorical values	hallucinated graph vocabulary or values
Direction validity	observed relationship directions	reversed or invalid traversals
Cypher analysis and normalization	structure extraction, rewrite safety, RETURN DISTINCT	unsafe rewrites or duplicate/noisy rows
Question constraints	quoted values use exact literals	fuzzy matching of exact user values
Execution	query runs and returns rows	syntactic validity without answerability
LLM judge	ambiguity, semantic alignment, schema use	executable but semantically weak examples

Table 1: PIPE-Cypher candidate acceptance gates.

ery run records accepted and rejected candidates for yield analysis. Retrieved few-shot examples replace graph-specific values with typed placeholders, preserving query structure while reducing tenant-value leakage and memorized entity reuse. A dependency-light value grounder adds typed prompt annotations for categorical values and reverse-bound entities, covering punctuation variants, possessives, plurals, synonyms, name partials, and small typos.

Inspired by Mind the Query’s prompt-setting analysis, the implementation exposes prompt profiles for schema-only, instructions-only, examples-only, examples-plus-instructions, and full governed PIPE-Cypher generation. These profiles are configured as ablation variants and must pass the same target-size and paper-readiness standards before any result is reported.

For LLM-judge review, the implementation slices schema context to the labels, relationship types, and properties used by the candidate query. This preserves validation against the full schema while keeping local 9B prompts within context limits on larger schemas.

The deterministic Cypher layer uses production-derived constraints: schema-only prompting, exact matching, relationship direction discipline, RETURN DISTINCT, reserved variable rejection, categorical values, required contextual return columns, fuzzy value annotations, placeholderized retrieval examples, and parser-aware rewrite boundaries. PIPE-Cypher records parser-style structure features and skips rewrites for risky constructs such as UNION, CALL, UNWIND, WHERE EXISTS, multiple WHERE clauses, or reserved variables.

We also estimate seed-template capacity before full generation. This check caught and fixed a scale blocker: reverse-binding execution was hard-capped to 10 rows, and no-slot negation/ranking seeds could not support full category targets. PIPE-Cypher now uses the configured binding limit and includes slotted negation and ranking seeds for both graphs.

For arbitrary enterprise onboarding, PIPE-Cypher includes a deployment template for a company’s own graph, read-only credentials, local model endpoint, schema introspection, privacy policy, dry run, scaled run, audit, and redacted export.

Graph	Target/category	Binding limit	Min seed capacity	All categories meet target
FinBench	250	300	300	Yes
SNB	125	200	200	Yes

Table 2: Built-in seed-template capacity after removing the reverse-binding 10-row cap and adding slotted negation/ranking seeds. Capacity is a theoretical scale-readiness check; final examples still must pass validation, execution, diversity, and judge gates.

Configurable value-sampling policies bound which low-cardinality graph values enter schema prompts, and redacted exports replace quoted literals, entity values, and string-valued result samples with stable placeholders for broader internal review. The ICIJ onboarding run further motivated schema-derived sparse-category templates: PIPE-Cypher now derives relationship-count, anti-join, and top-k templates from observed labels, relationship directions, and safe low-cardinality properties, then grounds slot values with outcome-aware reverse Cypher.

5 Experiments

Research questions. We evaluate PIPE-Cypher around four questions that matter for an industry benchmark generator: RQ1, can a local-model pipeline produce a balanced executable benchmark over live property graphs? RQ2, do Cypher-specific governance and grounding steps make generation reliable at scale? RQ3, does the resulting benchmark expose meaningful downstream Text2Cypher failures rather than merely checking syntax? RQ4, can the same machinery onboard a new public enterprise-style graph without hard-coding FinBench or SNB?

The generated benchmark contains 3,000 accepted examples: 2,000 from LDBC FinBench and 1,000 from LDBC SNB. Categories include simple retrieval, complex retrieval, simple aggregation, complex aggregation, boolean existence, negation/difference, path/temporal transaction, and ranking/top-k. Appendix Table 9 reports graph statistics, including the audited ICIJ third-graph onboarding workload.

We report three completed ablation suites over FinBench and SNB: an initial target-50 suite, a corrected target-100 suite, and a seed-17 target-50 repeat. Each suite contains 14 graph/setting cells over unconstrained local generation, reverse-only generation, validators+repair, no retrieval, no rewrite, no LLM judge, and full PIPE-Cypher. All paper-reported ablations have collection manifests, model IDs, code revisions, run summaries, and readiness

Setting	Value
Accepted examples	3,000
Primary graph	LDBC FinBench, 2,000 examples
Secondary graph	LDBC SNB, 1,000 examples
Categories	8 balanced categories, 375 each
Generation/judge model	Local Qwen3.5-9B
Execution backend	Neo4j Community, two live databases
Judge audit packet	80 sampled accepted/rejected pairs

Table 3: Full live experimental setup used for the generated benchmark artifact.

Graph	Candidates	Accepted	Acceptance	Categories at target
FinBench	3,404	2,000	0.588	8/8
SNB	1,373	1,000	0.728	8/8
Total	4,777	3,000	0.628	16/16

Table 4: Full live generation with local Qwen3.5-9B. Candidate counts include the initial sequential run and category-specific recovery top-ups.

audits. We additionally report a live ICIJ Offshore Leaks onboarding run as public finance/compliance evidence for arbitrary-schema deployment beyond LDBC.

Downstream Text2Cypher evaluation uses execution accuracy and answer-set precision/recall/F1 as primary correctness metrics. The benchmark evaluator also supports supplementary reference-based text metrics over serialized answer sets and query strings, including ROUGE, BLEU, METEOR, BERTScore, FrugalScore, cosine similarity, Jaro-Winkler similarity, and exact match; Appendix Table 24 lists the supported metrics and citations. We treat these as debugging and near-match diagnostics rather than substitutes for executable correctness. Appendix Table 12 and Figures 7–8 further separate workload-category balance from Cypher operator-strategy coverage.

6 Results

RQ1: executable benchmark generation. The full live run produced 3,000 accepted examples from 4,777 candidates using local Qwen3.5-9B for generation and judging. Category-specific recovery top-ups filled the only under-target categories from the initial sequential run. Every exported example passed read-only, syntax, schema, execution, non-empty result, and judge gates.

The accepted full-run records were exported into a benchmark package with stable example identifiers, train/dev/test splits, result samples, gate metadata, aggregate statistics, and a manifest hash for artifact tracking. The export contains 3,000 accepted examples: 2,000 FinBench, 1,000 SNB, and

Export artifact	Examples	FinBench	SNB	Train	Dev	Test
Live full benchmark	3,000	2,000	1,000	2,408	296	296

Table 5: Accepted live full benchmark package with stable IDs, gate metadata, result samples, statistics, and manifest hash 8b3c79a53a06b291a.

Artifact property	Value
Categories	8 balanced categories
FinBench/category	250
SNB/category	125
Difficulty split	1,521 easy / 1,479 medium
Unique labels / rel. types	7 / 9
Read/syntax/schema/exec/judge gates	3,000/3,000

Table 6: Distribution and gate summary for the exported full benchmark artifact.

375 accepted examples in every planned category.

Diversity and residual concentration. We report diversity as an audit signal, not as a slogan. The full export has perfect category balance and near-perfect difficulty balance, uses 1,115 unique grounded entity values, and exactly quotes grounded values in 82.6% of examples with entity bindings. The reported local-model run still has low query-signature diversity because seeded graph-grounded templates are doing substantial work. This supports the intended industry workflow: PIPE-Cypher can produce a balanced executable benchmark, and it also exposes the residual template and value concentration that a benchmark owner should monitor during refresh or fine-tuning data generation.

RQ2: governed generation. The full-run failure taxonomy is concentrated in diversity/duplicate controls during recovery and empty execution results; only 2/4,777 candidates were schema-invalid after deterministic Cypher governance. The scaled ablation suites should be read as a reliability study of the execution-grounded core rather than as a claim that every optional module increases yield. In the target-100 suite, every non-unconstrained graph/setting cell reached all eight category targets with 800 accepted examples per cell; the full PIPE-Cypher setting accepted 800/824 candidates on both FinBench and SNB. Figure 2 shows the contrast with unconstrained local generation, which produced no accepted examples under the same executable benchmark gates. Across the three evidence-ready suites, target-normalized coverage is 1.000 for every non-unconstrained cell, indicating that reverse grounding plus deterministic Cypher governance is a robust minimum viable

Split	Examples	Parse	Schema	Exec. success	Exec. acc.	Answer F1
Live full test	296	0.959	0.905	0.622	0.189	0.189

Table 7: Downstream Text2Cypher evaluation for local Qwen3.5-9B on the exported full benchmark test split.

core for enterprise benchmark generation.

Judge calibration. The released artifacts include an 80-row audit packet sampled from accepted and rejected full-run candidates, with a 40/40 judge accept/reject split and coverage over both graphs and all eight categories. A completed human audit shows 80.0% agreement, Cohen’s $\kappa = 0.60$, judge precision and specificity of 1.00, recall of 0.714, and no false accepts in the sample. The judge is therefore conservative: it protects accepted-example quality while rejecting some human-approved candidates and reducing yield. Human labels are used only for post-hoc calibration, not as a generation gate.

RQ3: downstream stress test. We ran downstream Text2Cypher evaluation on the exported full test split. Each model saw schema text and the natural-language question, generated Cypher, and was evaluated by live execution against the corresponding FinBench or SNB database. The purpose is not to present a single baseline as strong; it is to test whether PIPE-Cypher exposes failures in plausible Text2Cypher outputs. A 12-model local transfer suite spanning general, code-tuned, and Text2Cypher-tuned models confirms this role: zero-shot execution accuracy ranges from 0.000 to 0.291, while retrieval few-shot evaluation using PIPE-Cypher training examples ranges from 0.142 to 0.997 and improves 11/12 models (Appendix Table 8). Public Text2Cypher-style tuning therefore does not automatically transfer to enterprise-style graphs, but graph-specific examples make the benchmark immediately useful for retrieval, adaptation, and tenant-specific training data creation.

RQ4: third-graph onboarding. ICIJ onboarding reaches the same per-category target on a larger public finance/compliance graph: 800 accepted examples from 983 candidates over 2.0M nodes and 3.3M relationships, with 100 examples in every category. This result is important because it exercises schema-derived sparse-category templates and live reverse grounding on a graph whose schema is not one of the two LDBC study workloads.




Figure 2: Target-100 ablation yield on FinBench and SNB. Unconstrained local generation produces no accepted examples under executable benchmark gates, while every execution-grounded governed variant reaches all eight workload-category targets on both graphs. The takeaway is reliability of the grounded governance core, not that every optional module is needed for yield.

7 Industry Use

PIPE-Cypher is intended to be rerun when an enterprise graph changes. The generated artifact records the schema snapshot, graph profile, model identifier, validation gates, execution samples, judge scores, difficulty features, and source run for every example. Long-running jobs can be monitored from JSONL records and recovered with category-specific top-up runs that reject questions accepted in earlier runs. This turns benchmark maintenance into an operational workflow: refresh after schema changes, compare new models on the same held-out split, inspect failure modes by category, and export redacted review packets for governance teams.

The deployment path is designed for arbitrary enterprise schemas: teams provide read-only graph credentials, introspect labels/relationships/properties, choose a bounded categorical-value sampling policy, connect a local model endpoint, run a dry pass, scale generation, calibrate the judge, and export either raw internal artifacts or redacted review artifacts. This makes the FinBench/SNB/ICIJ study a reproducible evaluation setting rather than a hard-coded graph assumption.

8 Conclusion

PIPE-Cypher reframes Text2Cypher benchmarking as a private, repeatable enterprise workflow. The main lesson is that benchmark generation improves

when Cypher constraints become executable gates: reverse grounding makes questions answerable, deterministic governance catches unsafe or schema-invalid queries, execution exposes empty or brittle candidates, diversity diagnostics reveal concentration, and a calibrated local judge provides a conservative semantic filter. This produces a benchmark that is not only larger, but more useful: it is balanced, auditable, refreshable, and able to reveal downstream model failures that syntax-only evaluation would hide.

Limitations

Execution validity does not guarantee semantic correctness. The completed audit suggests a conservative judge with no false accepts in the labeled sample, but larger audits may reveal additional failure modes. Generated artifacts may contain private schema details or sampled values, so organizations should apply privacy redaction and normal data governance controls before broad sharing.

Ethics Statement

PIPE-Cypher is designed for private benchmark generation under local-model deployment constraints. Organizations using it should restrict benchmark artifacts to authorized users, review sampled values for sensitive content, and document whether judge calibration relied on human annotators.

References

- Satanjeev Banerjee and Alon Lavie. 2005. [METEOR: An automatic metric for MT evaluation with improved correlation with human judgments](#). In *Proceedings of the ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization*, pages 65–72, Ann Arbor, Michigan. Association for Computational Linguistics.
- Vashu Chauhan, Shobhit Raj, Shashank Mujumdar, Avirup Saha, and Anannay Jain. 2025. [Mind the query: A benchmark dataset towards Text2Cypher task](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: Industry Track*, pages 1890–1905, Suzhou, China. Association for Computational Linguistics.
- Matthew A. Jaro. 1989. [Advances in record-linkage methodology as applied to matching the 1985 census of tampa, florida](#). *Journal of the American Statistical Association*, 84(406):414–420.
- Moussa Kamal Eddine, Guokan Shang, Antoine Tixier, and Michalis Vazirgiannis. 2022. [FrugalScore: Learning cheaper, lighter and faster evaluation metrics for automatic text generation](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1305–1318, Dublin, Ireland. Association for Computational Linguistics.
- Fangyu Lei, Jixuan Chen, Yuxiao Ye, Ruisheng Cao, Dongchan Shin, Hongjin Su, Zhaoqing Suo, Hongcheng Gao, Wenjing Hu, Pengcheng Yin, Victor Zhong, Caiming Xiong, Ruoxi Sun, Qian Liu, Sida Wang, and Tao Yu. 2024. [Spider 2.0: Evaluating language models on real-world enterprise text-to-SQL workflows](#). *Preprint*, arXiv:2411.07763.
- Jinyang Li, Binyuan Hui, Ge Qu, Jiayi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Rongyu Cao, Ruiying Geng, Nan Huo, Xuanhe Zhou, Chenhao Ma, Guoliang Li, Kevin C. C. Chang, Fei Huang, Reynold Cheng, and Yongbin Li. 2023. [Can LLM already serve as a database interface? a BIG bench for large-scale database grounded text-to-SQLs](#). In *Advances in Neural Information Processing Systems*.
- Jiwei Li, Michel Galley, Chris Brockett, Jianfeng Gao, and Bill Dolan. 2016. [A diversity-promoting objective function for neural conversation models](#). In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 110–119, San Diego, California. Association for Computational Linguistics.
- Chin-Yew Lin. 2004. [ROUGE: A package for automatic evaluation of summaries](#). In *Text Summarization Branches Out*, pages 74–81, Barcelona, Spain. Association for Computational Linguistics.
- Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. 2008. [Introduction to Information Retrieval](#). Cambridge University Press.
- Makbule Gulcin Ozsoy, Leila Messallem, Jon Besga, and Gianandrea Minneci. 2025. [Text2Cypher: Bridging natural language and graph databases](#). In *Proceedings of the Workshop on Generative AI and Knowledge Graphs*, pages 100–108, Abu Dhabi, UAE. International Committee on Computational Linguistics.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. [Bleu: a method for automatic evaluation of machine translation](#). In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318, Philadelphia, Pennsylvania, USA. Association for Computational Linguistics.
- David P
üroja, Jack Waudby, Peter Boncz, and G
’abor Sz
’arnyas. 2023. [The LDBC social network benchmark interactive workload v2: A transactional graph query benchmark with deep delete operations](#). *Preprint*, arXiv:2307.04820.
- Shipeng Qi, Heng Lin, Zhihui Guo, G
’abor Sz
’arnyas, Bing Tong, Yan Zhou, Bin Yang, Jiansong Zhang, Zheng Wang, Youren Shen, Changyuan Wang, Parviz Peiravi, Henry Gabb, and Ben Steer. 2023. [The LDBC financial benchmark](#). *Preprint*, arXiv:2306.15975.
- Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. 2016. [SQuAD: 100,000+ questions for machine comprehension of text](#). In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392, Austin, Texas. Association for Computational Linguistics.
- Aman Tiwari, Shiva Krishna Reddy Malay, Vikas Yadav, Masoud Hashemi, and Sathwik Tejaswi Madhusudhan. 2025. [Auto-cypher: Improving LLMs on cypher generation via LLM-supervised generation-verification framework](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2: Short Papers*, pages 623–640, Albuquerque, New Mexico. Association for Computational Linguistics.
- William E. Winkler. 1990. [String comparator metrics and enhanced decision rules in the Fellegi-Sunter model of record linkage](#). In *Proceedings of the Section on Survey Research Methods*, pages 354–359. American Statistical Association.
- Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. 2020. [BERTScore: Evaluating text generation with BERT](#). In *International Conference on Learning Representations*.
- Xiuwen Zheng, Arun Kumar, and Amarnath Gupta. 2024. [Generating cross-model analytics workloads using LLMs](#). In *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management*, pages 4303–4307. ACM.

Zijie Zhong, Linqing Zhong, Zhaoze Sun, Qingyun Jin, Zengchang Qin, and Xiaofan Zhang. 2025. [SyntheT2C: Generating synthetic data for fine-tuning large language models on the Text2Cypher task](#). In *Proceedings of the 31st International Conference on Computational Linguistics*, pages 672–692, Abu Dhabi, UAE. Association for Computational Linguistics.

Yaoming Zhu, Sidi Lu, Lei Zheng, Jiaxian Guo, Weinan Zhang, Jun Wang, and Yong Yu. 2018. [Taxygen: A benchmarking platform for text generation models](#). In *The 41st International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 1097–1100. ACM.

A Additional Results

A.1 Downstream Stress-Test Interpretation

The downstream experiment is a benchmark-utility stress test. A useful enterprise benchmark should not merely reward syntactically valid Cypher; it should reveal when a model produces a query that parses, mentions plausible schema elements, and even executes, but answers the wrong operational question. The Qwen3.5-9B result has that shape: parse validity is 0.959 and schema validity is 0.905, while exact execution accuracy is only 0.189 on the full held-out split. This gap is evidence that PIPE-Cypher is measuring semantic graph-query competence rather than only query formatting.

The strategy-conditioned analysis makes the same point at operator level. Aggregation and single-hop queries account for most exact matches, whereas join-heavy, negation, ranking, and path-oriented examples expose wrong-direction, missing-context, execution-failure, and answer-mismatch errors. These zero-accuracy slices should not be read as a weakness of benchmark generation; they are the operational failure modes the benchmark is designed to surface. In an enterprise deployment, the same slices identify where retrieval examples, schema documentation, constrained decoding, or tenant-specific fine-tuning data are likely to matter.

We therefore ran the multi-model transfer stress test on the same fixed held-out split. The suite compares local general instruction, code-tuned, Cypher instruction, and Text2Cypher-finetuned models while keeping the schema prompt, evaluation backend, train/dev/test split, and execution metrics fixed. The result is consistent on the transfer claim: zero-shot execution accuracy is weak across model families, including public Text2Cypher-tuned checkpoints. Retrieval few-shot prompting with PIPE-

Cypher training examples usually closes much of the gap, but one LoRA run remains weak because its few-shot outputs are often syntactically malformed. This supports a practical industry reading of the benchmark: it is not only an evaluation artifact, but also a source of graph-specific examples for retrieval systems and future tenant-specific fine-tuning experiments, while still exposing model-family failures that require engineering attention.

B Claim–Evidence Map

This section maps the main paper claims to the strongest current evidence and to the remaining risks. This is intended as a reviewer-facing audit surface: claims with pending evidence remain marked as such rather than being folded into the main results.

1. **Claim 1.** PIPE-Cypher can generate a private, executable NL-to-Cypher benchmark over enterprise-style property graphs using local models.

Evidence. The full local Qwen3.5-9B run exported 3,000 accepted examples: 2,000 FinBench and 1,000 SNB, with 375 examples in every planned workload category and all accepted examples passing read-only, syntax, schema, execution, non-empty-result, and judge gates.

Key artifacts. benchmark export; manifest snapshot; paper table: benchmark export; paper table: full results

Status. Supported by live local-model evidence.

Risk. Generation quality may change with private enterprise schemas beyond FinBench/SNB.

2. **Claim 2.** Cypher-specific governance makes generated benchmark candidates safer and more auditable than prompt-only generation.

Evidence. The pipeline validates read-only safety, schema-visible labels, node and relationship properties, relationships, observed relationship direction, categorical values, contextual return columns, parser-style structure, execution success, and LLM-judge review. In the full run, only 2 of 4,777 candidates were schema-invalid after deterministic governance, and all 3,000 exported examples passed the deterministic and judge gates.

Model	Tuning	Zero acc.	Few-shot acc.	Δ	Few schema
Qwen3.5-9B	general instruction	0.189	0.997	0.807	1.000
Qwen2.5-Coder-7B-Instruct	code instruction	0.291	0.993	0.703	1.000
Gemma-2-9B-IT	general instruction	0.152	0.993	0.841	1.000
tomasonjo/text2cypher-demo-16bit	Text2Cypher fine-tuned	0.166	0.976	0.811	1.000
neo4j/Gemma-3-4B Text2Cypher	Text2Cypher fine-tuned	0.000	0.912	0.912	0.997
Azzedde/llama3.1-8b-text2cypher	Text2Cypher fine-tuned	0.155	0.976	0.821	0.997
neo4j/Gemma-2-9B Text2Cypher LoRA	Text2Cypher LoRA	0.199	0.993	0.794	0.997
aigents/Llama-3.1-8B Cypher LoRA	Cypher LoRA	0.074	0.730	0.655	0.912
aigents/Llama-3.1-8B Cypher mixed LoRA	Cypher mixed LoRA	0.111	0.983	0.872	0.990
ragraph-ai/stable-cypher-instruct-3b	Cypher instruction	0.155	0.838	0.682	0.976
projectwilson/Llama-3.1-8B Text2Cypher LoRA	Text2Cypher LoRA	0.169	0.142	-0.027	0.990
Saiprasanth15/Llama-3.1-8B Text2Cypher LoRA	Text2Cypher LoRA	0.000	0.730	0.730	0.993

Table 8: Local-model downstream Text2Cypher transfer stress test on the same full held-out PIPE-Cypher test split. The table reports only runs with completed zero-shot and retrieval few-shot summaries.

Graph	Nodes	Relationships	Labels	Rel. types	Paper status
FinBench	10,006	57,622	5	9	reported
SNB	34,735	70,842	14	15	reported
ICIJ Offshore Leaks	2,016,523	3,339,267	5	14	onboarding audit

Table 9: Study graph statistics. ICIJ Offshore Leaks is used as a public third-graph onboarding audit for arbitrary finance/compliance schemas beyond the two LDBC study workloads.

ICIJ onboarding property	Value
Graph nodes / relationships	2,016,523 / 3,339,267
Labels / relationship types	5 / 14
Generated / accepted	983 / 800
Acceptance rate	0.814
Categories at target	8/8
Paper audit	ready
Sparse schema-derived accepts	complex agg. 97, negation 28, ranking 98

Table 10: ICIJ Offshore Leaks third-graph onboarding audit. The public finance/compliance graph tests arbitrary-schema generation beyond the two LDBC study workloads; raw values remain outside the paper artifacts.

Key artifacts. code: validator.py; code: cypher_parser.py; code: question_constraints.py; snapshot: failure taxonomy; +1 more

Status. Supported by code, tests, and full-run failure taxonomy.

Risk. Semantic correctness still depends on execution evidence and judge review; the completed 80-row audit is a calibration sample rather than an exhaustive proof.

- Claim 3.** Reverse grounding and structured workload seeds are designed to improve reliability under local-model constraints.

Evidence. Three FinBench/SNB ablation suites are complete and evidence-ready: the original target-50 suite, the corrected target-100 suite, and a seed-17 target-50 repeat.

Metric	Value
Question Distinct-1	0.041
Question Distinct-2	0.088
Question self-BLEU-2 (sampled)	0.826
Unique query-signature ratio	0.010
Top query-signature share	0.083
Category normalized entropy	1.000
Graph-category normalized entropy	0.980
Difficulty normalized entropy	1.000
Label coverage	0.412
Relationship-type coverage	0.375
Property-name coverage	0.407
Unique grounded-value ratio	0.373
Grounded values exactly quoted	0.826
Aggregation / negation / ordering rates	0.500 / 0.125 / 0.125

Table 11: Diversity diagnostics for the full exported benchmark. Distinct-n follows text-generation usage; self-BLEU is lower when questions are less redundant; grounded-value metrics are aggregate-only and do not list raw values.

Primary strategy	Examples	Share	Rel. patterns	Return arity	Downstream exec. acc.
Aggregation	1125	0.375	1.42	1.00	0.396
Join-heavy	375	0.125	2.33	2.33	0.000
Negation	375	0.125	2.08	3.00	0.000
Order/rank	375	0.125	1.99	3.34	0.000
Path	375	0.125	1.37	3.00	0.000
Single hop	375	0.125	1.00	2.33	0.324

Table 12: Cypher strategy diagnostics over the full 3,000-example export. Strategy tags are derived from generated Cypher structure rather than from category labels; downstream execution accuracy is reported on the full held-out test split when a strategy appears there.

Each suite completed 14/14 graph/variant cells, reached every category target for all non-empty/non-unconstrained runs, and passed paper-readiness audits with logs, model IDs, code revisions, run summaries, collection manifests, and gate-rate availability. The target-100 suite is the detailed appendix ablation table/figure, while the three-suite comparison reports target-normalized coverage, acceptance variation, execution rates, and judge rates for stability.

Failure bucket	Count	Share of rejected
Diversity/duplicate control	1,335	0.751
Empty result	374	0.210
Judge semantic reject	66	0.037
Schema invalid	2	0.001

Table 13: Failure taxonomy over full-run generation candidates before benchmark export. Accepted examples are excluded from the bucket shares.

PIPE-Cypher category	Closest MTQ category	Added enterprise distinction
Simple retrieval	SR	Direct node/edge lookup with exact filters.
Complex retrieval	CR	Multi-hop or multi-pattern retrieval.
Simple aggregation	SA	Single aggregation such as count/min/max/average.
Complex aggregation	CA	Grouped or multi-stage aggregation over graph neighborhoods.
Boolean existence	EQ	Precise yes/no or existence answer.
Negation/difference	CR	Absence, anti-join, or difference query.
Path/temporal transaction	CR/CA	Temporal or path-oriented transaction neighborhood.
Ranking/top-k	SA/CA	Ordered top-k query with explicit limit.

Table 14: Category crosswalk to Mind the Query. PIPE-Cypher keeps the familiar retrieval/aggregation/evaluation-query structure while adding enterprise workloads such as negation, temporal paths, and ranking.

Gate or ledger	Count	Denominator
Export accepted	3,000	3,000
Read-only safety	3,000	3,000
Syntax validity	3,000	3,000
Schema/value validity	3,000	3,000
Execution success	3,000	3,000
LLM judge pass	3,000	3,000
Rejected candidates logged	1,777	4,777

Table 15: PIPE-Cypher validation cascade for the full export and its logged candidate ledger. Unlike Mind the Query, human review is calibration-only.

Profile	Added constraint	Intended evidence
Schema-only	Use only visible schema.	Baseline for schema-grounding prompting.
Instructions	Exact values, datatype rules, no nested aggregations.	Targets MTQ-style prompt refinements.
Examples	Placeholderized retrieved NL-Cypher pairs.	Tests whether examples help without extra governance.
Examples + instructions	Few-shot plus explicit rules.	Closest controlled analogue to MTQ Table 5.
Full governed	Production-derived Cypher hints, AST-safe rewrites, judge gate.	PIPE-Cypher production setting.

Table 16: Prompt-profile plan inspired by Mind the Query’s prompt-variant analysis. Only completed, audited target-50-or-larger results should be promoted into results tables.

Key artifacts. script: run live ablation suite; script: monitor remote ablation suite; script: monitor remote ablation queue; script: collect remote ablation suite; +20 more

Status. Supported by audited target-100 evidence and three-suite target-size/repeated-seed comparison.

Risk. The unconstrained local generation baseline produced no accepted examples under the strict executable gates, so comparisons against it should be framed as a gate-stress baseline rather than a tuned generation system.

- Claim 4.** The benchmark controls category and difficulty balance while exposing residual diversity limits.

Evidence. The full export has perfect category balance, planned 2:1 FinBench/SNB graph mix, and a near-even easy/medium split. Diversity diagnostics report Distinct-n, sampled self-BLEU-2, query-signature diversity, normalized entropy, schema coverage, structural feature rates, and aggregate value-grounding diagnostics. The full export uses 1,115 unique grounded entity values and exactly quotes grounded values in 82.6% of examples with entity bindings, making template and value concentration visible instead of hidden.

Key artifacts. snapshot: diversity report; paper table: diversity; paper figure: diversity diagnostics; paper figure: full export distribution

Status. Supported by full-export statistics and diversity diagnostics.

Risk. Low query-signature diversity in the reported local-model run remains a limitation and ablation target.

- Claim 5.** The generated benchmark is useful for downstream Text2Cypher evaluation.

Evidence. A zero-shot local Qwen3.5-9B downstream evaluation over the 296-example full test split reports parse validity, schema validity, execution success, execution accuracy, answer F1, per-category breakdowns, fixed-seed bootstrap uncertainty intervals, and a row-level error taxonomy. The model reaches 0.189 execution accuracy with a 95% bootstrap interval of [0.145, 0.233] and 0.622 execution success with interval [0.564, 0.676]. The error taxonomy classifies 240 incorrect rows into answer mismatches, execution failures, schema invalid predictions, and parse-invalid predictions, showing the benchmark can discriminate easy aggregation from harder operational categories and can explain model failure modes.

Key artifacts. downstream summary; snapshot: downstream uncertainty; snapshot: downstream error report; research note: downstream evaluation; +6 more

Status. Supported by live downstream evaluation against FinBench/SNB databases, with

Dimension	Mind the Query	PIPE-Cypher
Generation review gate	Manual logical review	Deterministic gates + local LLM judge
Human effort	Reported 1,400 person-hours	80-row post-hoc judge calibration audit
Private values	Public benchmark values	Configurable sampling and redacted export
Refresh	Static dataset release	Rerunnable private benchmark factory
Model endpoint	Gemini for generation	Local Qwen3.5-9B endpoint

Table 17: Industry deployment contrast with Mind the Query. PIPE-Cypher focuses on private refreshable benchmark generation rather than a one-time public dataset.

Setting	Graph	Records	Accepted	Acceptance	Categories at target
Unconstrained LLM	FinBench	0	0	0.000	0/8
Unconstrained LLM	SNB	0	0	0.000	0/8
Reverse-only	FinBench	815	800	0.982	8/8
Reverse-only	SNB	820	800	0.976	8/8
Validators+repair	FinBench	833	800	0.960	8/8
Validators+repair	SNB	824	800	0.971	8/8
No retrieval	FinBench	819	800	0.977	8/8
No retrieval	SNB	809	800	0.989	8/8
No rewrite	FinBench	826	800	0.969	8/8
No rewrite	SNB	834	800	0.959	8/8
No LLM judge	FinBench	812	800	0.985	8/8
No LLM judge	SNB	828	800	0.966	8/8
Full PIPE-Cypher	FinBench	824	800	0.971	8/8
Full PIPE-Cypher	SNB	824	800	0.971	8/8

Table 18: Live target-100 ablation evidence with local Qwen3.5-9B. Each graph run targets 100 accepted examples per category.

bootstrap uncertainty computed from row-level outputs.

Risk. Only one downstream model has been evaluated so far.

- Claim 6.** Automated LLM-judge review can replace human-in-the-loop generation gates while still exposing judge risk.

Evidence. The pipeline uses deterministic validation plus LLM-judge review as the generation gate. A post-hoc 80-row judge audit packet preserves 40 judge accepts, 40 judge rejects, both graphs, and all eight categories. The completed human audit reports 80.0% agreement, Cohen’s kappa 0.60, balanced accuracy 0.857, judge precision 1.00, judge specificity 1.00, judge recall 0.714, zero false accepts, and 16 false rejects. The analyzer requires complete labels for paper-facing calibration, and the tracked summary avoids committing raw value-bearing audit rows.

Key artifacts. code: judge.py; code: calibration.py; code: judge_audit_packet.py; script: create judge annotation sheets; +6 more

Status. Supported by completed human audit summary.

Risk. The judge appears conservative on this sample; larger audits or different enterprise schemas may reveal false accepts.

Setting	Graph	Read-only	Syntax	Schema	Exec.	Judge
Unconstrained LLM	FinBench	0.000	0.000	0.000	0.000	0.000
Unconstrained LLM	SNB	0.000	0.000	0.000	0.000	0.000
Reverse-only	FinBench	1.000	1.000	0.982	0.982	0.982
Reverse-only	SNB	1.000	1.000	0.998	0.998	0.976
Validators+repair	FinBench	1.000	1.000	1.000	1.000	0.960
Validators+repair	SNB	1.000	1.000	0.998	0.998	0.971
No retrieval	FinBench	1.000	1.000	1.000	1.000	0.977
No retrieval	SNB	1.000	1.000	0.998	0.998	0.989
No rewrite	FinBench	1.000	1.000	1.000	1.000	0.969
No rewrite	SNB	1.000	1.000	0.998	0.998	0.959
No LLM judge	FinBench	1.000	1.000	1.000	1.000	0.985
No LLM judge	SNB	1.000	1.000	0.998	0.998	0.966
Full PIPE-Cypher	FinBench	1.000	1.000	1.000	1.000	0.971
Full PIPE-Cypher	SNB	1.000	1.000	0.998	0.998	0.971

Table 19: Quality-gate rates for the live target-100 ablation suite. Rates are computed over all generated records in each graph/setting.

- Claim 7.** PIPE-Cypher supports arbitrary enterprise schema onboarding with configurable privacy redaction and value-sampling policies.

Evidence. The repo includes an enterprise deployment guide, an enterprise template config, privacy config fields, schema introspection that reads categorical-value thresholds, maximum sampled value length, and omitted-property patterns from config, deterministic redaction helpers, and a redacted benchmark export CLI. The redaction policy replaces quoted Cypher literals, question mentions, entity values, and string-valued result samples with stable placeholders while preserving query structure and gate metadata. The repo now also includes an ICIJ Offshore Leaks onboarding profile, config, download/load helpers, graph plan, schema-derived sparse-category templates, and a completed sanitized target-100 onboarding audit on a public finance/compliance graph beyond the LDBC study workloads.

Key artifacts. research note: enterprise deployment guide; research note: icij offshore-leaks graph plan; enterprise_template.yaml; schema_icij_offshoreleaks.json; +21 more

Status. Supported by docs, config, code, deterministic tests, a concrete third-graph onboarding plan, and a completed sanitized ICIJ target-100 live onboarding run with 800/983 accepted examples and 8/8 categories at target.

Risk. ICIJ is public rather than private/anonymized enterprise data. The strongest industry validation would still come from a real tenant graph under the same audit protocol.

Setting	Graph	Suites	Target cov.	Acceptance	Cat. target	Exec.	Judge
No LLM judge	FinBench	3/3	1.000	0.994±0.008	1.000	1.000	0.994
No retrieval	FinBench	3/3	1.000	0.967±0.031	1.000	1.000	0.967
No rewrite	FinBench	3/3	1.000	0.969±0.023	1.000	1.000	0.969
Full PIPE-Cypher	FinBench	3/3	1.000	0.975±0.016	1.000	1.000	0.975
Reverse-only	FinBench	3/3	1.000	0.994±0.011	1.000	0.994	0.994
Unconstrained LLM	FinBench	3/3	0.000	0.000	0.000	0.000	0.000
Validators+repair	FinBench	3/3	1.000	0.987±0.023	1.000	1.000	0.987
No LLM judge	SNB	3/3	1.000	0.982±0.015	1.000	0.996	0.982
No retrieval	SNB	3/3	1.000	0.993±0.004	1.000	0.996	0.993
No rewrite	SNB	3/3	1.000	0.983±0.021	1.000	0.996	0.983
Full PIPE-Cypher	SNB	3/3	1.000	0.987±0.014	1.000	0.996	0.987
Reverse-only	SNB	3/3	1.000	0.989±0.011	1.000	0.996	0.989
Unconstrained LLM	SNB	3/3	0.000	0.000	0.000	0.000	0.000
Validators+repair	SNB	3/3	1.000	0.987±0.014	1.000	0.996	0.987

Table 20: Target-size and repeated-seed ablation sensitivity. Target coverage normalizes accepted examples by each suite’s planned graph/category target, so target-50 and target-100 suites can be compared without treating larger raw counts as quality gains.

Audit packet property	Value
Rows	80
Judge accept / reject	40 / 40
FinBench / SNB rows	48 / 32
Easy / medium rows	26 / 54
Labeled rows	80
Agreement / Cohen’s κ	0.800 / 0.600
Judge precision / recall	1.000 / 0.714
Judge specificity	1.000
False accepts / false rejects	0 / 16

Table 21: Post-hoc judge calibration from the completed 80-row human audit. Labels are used only to calibrate the automated judge, not as a generation gate.

Metric	Point	95% CI	SE	N
Parse valid	0.959	[0.936, 0.980]	0.012	296
Schema valid	0.905	[0.872, 0.939]	0.017	296
Execution success	0.622	[0.564, 0.676]	0.028	296
Execution accuracy	0.189	[0.145, 0.233]	0.022	296
Answer F1	0.189	[0.149, 0.236]	0.023	296

Table 22: Bootstrap uncertainty for downstream Text2Cypher evaluation on the full exported test split. Intervals use row-level resampling with a fixed seed and are intended for appendix reporting.

Downstream outcome	Count	Share of incorrect
Answer mismatch	128	0.533
Execution failed	72	0.300
Schema invalid	28	0.117
Parse invalid	12	0.050

Table 23: Downstream Text2Cypher failure taxonomy for local Qwen3.5-9B on the full exported test split. Shares exclude exact-answer matches.

C Prompt Contracts

PIPE-Cypher treats prompts as versioned implementation artifacts. The list summarizes the prompt contracts used for generation, repair, judging, and downstream evaluation; hashes fingerprint the full prompt constants in the codebase.

1. **Template generation.** Stage: Workload proposal. SHA-256: 10b25b.

Contract. Schema-only labels, relationships, properties, and categorical values; Realistic enterprise analyst wording; At most two typed slots and JSON-only output

2. **Reverse binding.** Stage: Graph grounding. SHA-256: 48e949.

Contract. Read-only MATCH/WHERE/RETURN DISTINCT/LIMIT only; Slot variables named

exactly as requested; Forward relationship directions from the schema

3. **Cypher generation.** Stage: Candidate query. SHA-256: 61c557.

Contract. Only schema-visible constructs and observed directions; RETURN DISTINCT for set returns and exact equality for quoted values; Context columns, categorical hints, placeholderized retrieval, and no writes

4. **Repair.** Stage: Validation feedback. SHA-256: 56c419.

Contract. Preserve question intent while fixing validation or execution issues; Keep query

Metric	PIPE-Cypher support	Primary citation
ROUGE-1/2/L	Reference overlap over serialized answer sets and query strings	(Lin, 2004)
BLEU	Sentence-level reference overlap with brevity penalty	(Papineni et al., 2002)
METEOR	Exact-token METEOR-style precision/recall with fragmentation penalty	(Banerjee and Lavie, 2005)
BERTScore	Optional contextual-embedding similarity when installed	(Zhang et al., 2020)
FrugalScore	Optional efficient learned NLG metric when installed	(Kamal Eddine et al., 2022)
Cosine	Token-count vector cosine similarity	(Manning et al., 2008)
Jaro-Winkler	Character-level string similarity for near-match diagnostics	(Jaro, 1989; Winkler, 1990)
Exact match	Raw and normalized string exact match	(Rajpurkar et al., 2016)

Table 24: Supplementary reference-based text metrics supported by the PIPE-Cypher evaluator. These metrics are useful for debugging answer rendering, paraphrase sensitivity, and near-match behavior, but they do not replace execution accuracy or answer-set F1 for Text2Cypher correctness. BERTScore and FrugalScore are optional integrations because they require additional metric/model packages.

841	read-only and schema-grounded; Return only	<i>Relationships:</i> TRANSFER_TO.	877
842	corrected Cypher	<i>Gates:</i> RO/Syn/Schema/Exec/Judge.	878
843	5. LLM judge. Stage: Quality gate. SHA-256:	<i>Result sample:</i> {HasOutgoingTransfer: True};	879
844	421c7b.	observed rows: 1	880
845	<i>Contract.</i> Inputs include question, Cypher,	2. FinBench / Complex Aggregation / medium.	881
846	schema slice, execution rows, and validation	<i>Question:</i> What is the total transferred amount	882
847	summary; Strict JSON scores for ambiguity,	from accounts owned by person 'Gwar'?	883
848	semantic alignment, schema use, and diffi-		
849	culty; Categorical values constrain query liter-	<pre>MATCH (p:Person {personName: 'Gwar'}) -[:OWN_ACCOUNT]-> (src:Account)-[t:TRANSFER_TO]-> (:Account) RETURN DISTINCT SUM(t.amount) AS TotalTransferredAmount</pre>	884
850	als, not observed result-row values; Pass only		885
851	useful, unambiguous enterprise benchmark		886
852	examples		887
853	6. Downstream Text2Cypher. Stage: Model	<i>Structure:</i> join_heavy, aggregation.	888
854	evaluation. SHA-256: 4c07ff.	<i>Relationships:</i> OWN_ACCOUNT, TRANS-	891
855	<i>Contract.</i> Read-only Cypher only; Schema-	FER_TO.	892
856	visible constructs and exact direction preser-	<i>Gates:</i> RO/Syn/Schema/Exec/Judge.	893
857	vation; RETURN DISTINCT, count/ranking	<i>Result sample:</i> {TotalTransferredAmount:	894
858	rules, and no explanations	14972318.41}; observed rows: 1	895
859	D Representative Accepted Examples	3. FinBench / Complex Retrieval / easy. Ques-	896
860	The examples below are selected in stable identifier	<i>tion:</i> Which accounts received transfers from	897
861	order from the tracked full-export snapshot, one per	accounts owned by person 'Zof'?	898
862	graph/category cell when available. They show the		899
863	natural-language question, accepted Cypher, struc-	<pre>MATCH (p:Person {personName: 'Zof'}) -[:OWN_ACCOUNT]-> (src:Account) -[:TRANSFER_TO]-> (dst:Account) RETURN DISTINCT dst.accountId AS AccountId, dst.accountType AS AccountType, dst.isBlocked AS IsBlocked LIMIT 300</pre>	900
864	tural tags, gate status, and a bounded execution-	<i>Structure:</i> join_heavy, bounded_result.	901
865	result sample.	<i>Relationships:</i> OWN_ACCOUNT, TRANS-	902
866	1. FinBench / Boolean Existence / easy. Ques-	FER_TO.	903
867	<i>tion:</i> Does account '187743809466009406'	<i>Gates:</i> RO/Syn/Schema/Exec/Judge.	904
868	have any outgoing transfer?	<i>Result sample:</i> {AccountId:	905
869	<pre>MATCH (src:Account {accountId: '187743809466009406'}) -[:TRANSFER_TO]-> (:Account) RETURN DISTINCT COUNT(DISTINCT src) > 0 AS HasOutgoingTransfer</pre>	4687402787162554854, AccountType:	906
870		credit card, IsBlocked: False}; observed rows:	907
871		3	908
872	<i>Structure:</i> single_hop, aggregation.		909
873			910
874			911
875			912
876			913
			914
			915
			916

4. FinBench / Negation Difference / medium.

Question: Which accounts owned by person 'Kant' have not sent any transfers?

```
MATCH (p:Person {personName:
'Kant'})
-[:OWN_ACCOUNT]-> (a:Account)
WHERE NOT (a) -[:TRANSFER_TO]->
(:Account)
RETURN DISTINCT a.accountId AS
AccountId, a.accountType AS
AccountType,
a.isBlocked AS IsBlocked
LIMIT 300
```

Structure: join_heavy, negation, bounded_result.

Relationships: OWN_ACCOUNT, TRANSFER_TO.

Gates: RO/Syn/Schema/Exec/Judge.

Result sample: {AccountId: 4739757132830738341, AccountType: merchant account, IsBlocked: False}; observed rows: 1

5. FinBench / Path Temporal / medium. *Ques-*

tion: Which accounts can receive money within two transfer hops from accounts owned by person 'Sossamon'?

```
MATCH (p:Person {personName:
'Sossamon'}) -[:OWN_ACCOUNT]->
(src:Account) -[:TRANSFER_TO*1..2]->
(dst:Account)
RETURN DISTINCT dst.accountId AS
AccountId, dst.accountType AS
AccountType, dst.isBlocked AS
IsBlocked
LIMIT 300
```

Structure: join_heavy, path, bounded_result.

Relationships: OWN_ACCOUNT, TRANSFER_TO.

Gates: RO/Syn/Schema/Exec/Judge.

Result sample: {AccountId: 4687402787162554854, AccountType: credit card, IsBlocked: False}; observed rows: 4

6. FinBench / Ranking Topk / medium. *Ques-*

tion: For accounts owned by person 'Barry', which account sent the highest total transfer amount?

```
MATCH (p:Person {personName:
'Barry'})
-[:OWN_ACCOUNT]->
(src:Account) -[:TRANSFER_TO]->
(:Account)
WITH src, SUM(t.amount) AS
totalAmount
RETURN DISTINCT src.accountId AS
AccountId, src.accountType AS
AccountType, src.isBlocked AS
IsBlocked,
totalAmount
```

```
ORDER BY totalAmount DESC
LIMIT 1
```

Structure: join_heavy, aggregation, order_rank, bounded_result.

Relationships: OWN_ACCOUNT, TRANSFER_TO.

Gates: RO/Syn/Schema/Exec/Judge.

Result sample: {AccountId: 4732438783436260772, AccountType: internet account, IsBlocked: False}; observed rows: 1

7. FinBench / Simple Aggregation / easy. *Ques-*

tion: How many accounts are owned by person 'Kaewsuktae'?

```
MATCH (p:Person {personName:
'Kaewsuktae'}) -[:OWN_ACCOUNT]->
(a:Account)
RETURN DISTINCT COUNT(DISTINCT a) AS
AccountCount
```

Structure: single_hop, aggregation.

Relationships: OWN_ACCOUNT.

Gates: RO/Syn/Schema/Exec/Judge.

Result sample: {AccountCount: 1}; observed rows: 1

8. FinBench / Simple Retrieval / easy. *Ques-*

tion: Which accounts are owned by person 'Barry'?

```
MATCH (p:Person {personName:
'Barry'})
-[:OWN_ACCOUNT]-> (a:Account)
RETURN DISTINCT a.accountId AS
AccountId, a.accountType AS
AccountType,
a.isBlocked AS IsBlocked
LIMIT 300
```

Structure: single_hop, bounded_result.

Relationships: OWN_ACCOUNT.

Gates: RO/Syn/Schema/Exec/Judge.

Result sample: {AccountId: 4732438783436260772, AccountType: internet account, IsBlocked: False}; observed rows: 1

9. SNB / Boolean Existence / medium. *Ques-*

tion: Does person with id 6597069766828 like any post?

```
MATCH (p:Person {id:
6597069766828})
OPTIONAL MATCH (p) -[:LIKES]->
(post:Post)
RETURN DISTINCT COUNT(post) > 0 AS
LikesAnyPost
```

1027	Structure: single_hop, aggregation, optional.	Gates: RO/Syn/Schema/Exec/Judge.	1080
1028	Relationships: LIKES.	Result sample: {FirstName: K., LastName: Sen, PersonId: 94}; observed rows: 1	1081
1029	Gates: RO/Syn/Schema/Exec/Judge.		1082
1030	Result sample: {LikesAnyPost: True}; observed rows: 1	13. SNB / Path Temporal / easy. Question: Which people are within two knows hops of person with id 4398046511136?	1083
1031			1084
1032	10. SNB / Complex Aggregation / medium. Question: How many distinct posts are in forums joined by person with id 6597069766845?		1085
1033		<pre>MATCH (src:Person {id: 4398046511136}) -[:KNOWS*1..2]-> (dst:Person) RETURN DISTINCT dst.id AS PersonId, dst.firstName AS FirstName, dst.lastName AS LastName LIMIT 200</pre>	1086
1034			1087
1035			1088
1036	<pre>MATCH (forum:Forum) -[:HAS_MEMBER]-> (p:Person {id: 6597069766845}) MATCH (forum) -[:CONTAINER_OF]-> (post:Post) RETURN DISTINCT COUNT(DISTINCT post) AS JoinedForumPostCount</pre>		1089
1037			1090
1038			1091
1039			1092
1040			1093
1041		Structure: single_hop, path, bounded_result.	1094
1042		Relationships: KNOWS.	1095
1043	Structure: join_heavy, aggregation.	Gates: RO/Syn/Schema/Exec/Judge.	1096
1044	Relationships: CONTAINER_OF,	Result sample: {FirstName: Rafael, LastName: Fernández, PersonId: 4398046511333}; observed rows: 25	1097
1045	HAS_MEMBER.		1098
1046	Gates: RO/Syn/Schema/Exec/Judge.		1099
1047	Result sample: {JoinedForumPostCount: 1}; observed rows: 1	14. SNB / Ranking Topk / medium. Question: Among forums tagged 'James_K._Polk', which forum has the most members?	1100
1048			1101
1049	11. SNB / Complex Retrieval / easy. Question: Which people are members of forums containing posts tagged 'Manuel_Noriega'?		1102
1050		<pre>MATCH (forum:Forum) -[:HAS_TAG]-> (tag:Tag {name: 'James_K._Polk'}) MATCH (forum) -[:HAS_MEMBER]-> (person:Person) WITH forum, COUNT(DISTINCT person) AS memberCount RETURN DISTINCT forum.title AS ForumTitle, memberCount ORDER BY memberCount DESC LIMIT 1</pre>	1103
1051			1104
1052	<pre>MATCH (forum:Forum) -[:HAS_MEMBER]-> (p:Person), (forum) -[:CONTAINER_OF]-> (post:Post) -[:HAS_TAG]-> (tag:Tag {name: 'Manuel_Noriega'}) RETURN DISTINCT p.id AS PersonId LIMIT 200</pre>		1105
1053			1106
1054			1107
1055			1108
1056			1109
1057			1110
1058			1111
1059	Structure: join_heavy, bounded_result.	Structure: join_heavy, aggregation, order_rank, bounded_result.	1114
1060	Relationships: CONTAINER_OF,	Relationships: HAS_MEMBER, HAS_TAG.	1115
1061	HAS_MEMBER, HAS_TAG.		1116
1062	Gates: RO/Syn/Schema/Exec/Judge.	Gates: RO/Syn/Schema/Exec/Judge.	1117
1063	Result sample: {PersonId: 4398046511220}; observed rows: 19	Result sample: {ForumTitle: Wall of Javed Khan, memberCount: 33}; observed rows: 1	1118
1064			1119
1065	12. SNB / Negation Difference / medium. Question: Which members of forum 'Wall of Celso Oliveira' have not liked any post?	15. SNB / Simple Aggregation / easy. Question: How many posts are tagged 'Vietnam'?	1120
1066			1121
1067		<pre>MATCH (post:Post) -[:HAS_TAG]-> (tag:Tag {name: 'Vietnam'}) RETURN DISTINCT COUNT(DISTINCT post) AS PostCount</pre>	1122
1068	<pre>MATCH (forum:Forum {title: 'Wall of Celso Oliveira'}) -[:HAS_MEMBER]-> (p:Person) WHERE NOT (p) -[:LIKES]-> (:Post) RETURN DISTINCT p.id AS PersonId, p.firstName AS FirstName, p.lastName AS LastName LIMIT 200</pre>		1123
1069			1124
1070			1125
1071			1126
1072			1127
1073		Structure: single_hop, aggregation.	1128
1074		Relationships: HAS_TAG.	1129
1075		Gates: RO/Syn/Schema/Exec/Judge.	1130
1076		Result sample: {PostCount: 1}; observed rows: 1	1131
1077	Structure: join_heavy, negation, bounded_result.		1132
1078	Relationships: HAS_MEMBER, LIKES.		1133




Figure 3: Gate-rate heatmap for the audited target-100 ablation suite. Deterministic read-only and syntax gates are saturated; execution and judge rates expose the remaining quality differences across graph and pipeline variants.




Figure 4: Three-suite ablation stability over the original target-50 suite, corrected target-100 suite, and seed-17 target-50 repeat. Target-normalized coverage stays at 1.000 for all non-unconstrained cells.

16. **SNB / Simple Retrieval / easy.** *Question:* Which post IDs did person with id 4398046511124 like?

```
MATCH (p:Person {id:
4398046511124})
-[:LIKES]-> (post:Post)
RETURN DISTINCT post.id AS PostId
LIMIT 200
```

Structure: single_hop, bounded_result.

Relationships: LIKES.

Gates: RO/Syn/Schema/Exec/Judge.

Result sample: {PostId: 343597385744}; observed rows: 5




Figure 5: ICIJ Offshore Leaks onboarding audit. Schema-derived sparse-category templates recover balanced target-100 coverage on a public finance/compliance graph beyond the two LDDB workloads.




Figure 6: Diversity diagnostics for the full 3,000-example export. Category, graph-category, and difficulty balance are strong by construction; schema coverage and query-signature diversity expose remaining concentration from the seeded-template run.




Figure 7: Cypher strategy coverage by workload category. Category balancing does not by itself guarantee operator coverage; the strategy matrix shows where the benchmark exercises single-hop retrieval, joins, aggregation, ordering, negation, path patterns, optional matches, and bounded-result queries.




Figure 8: Downstream outcomes by gold Cypher strategy on the full held-out test split. The local Qwen3.5-9B baseline fails differently across strategies: aggregation mixes exact matches with execution failures, while join-heavy, negation, path, and ranking examples are dominated by semantically wrong executable queries or execution failures.




Figure 9: Full 3,000-example export distribution. The benchmark preserves the planned 2:1 FinBench/SNB graph mix while balancing all eight Cypher workload categories and maintaining a near-even easy/medium difficulty split.




Figure 10: Bootstrap uncertainty for downstream Qwen3.5-9B Text2Cypher evaluation by workload category. Error bars show 95% row-level bootstrap intervals over the full exported test split.




Figure 11: Downstream Text2Cypher failure taxonomy for the local Qwen3.5-9B baseline on the full test split. Exact matches are excluded so the chart exposes whether misses come from invalid Cypher, failed execution, or semantically wrong executable queries.




Figure 12: Full-run rejection taxonomy before benchmark export. Most rejected candidates come from duplicate/diversity control during category recovery and from empty execution results; schema-invalid Cypher is rare after deterministic governance.