

PIPE-Cypher: Automatic Enterprise Benchmark Generation for Text-to-Cypher Systems

Anonymous ACL submission

Abstract

Enterprises increasingly use property graphs and Cypher to analyze fraud, risk, identity, customer, and operational data. Public Text2Cypher benchmarks rarely reflect private schemas, terminology, access patterns, or difficulty distributions. We present PIPE-Cypher, a local-model pipeline for generating organization-specific natural-language-to-Cypher benchmarks. PIPE-Cypher combines schema introspection, reverse-query grounding, constrained Cypher generation, deterministic read-only and schema validation, execution feedback, repair, diversity controls, and LLM-judge review. The system is designed for industry settings where benchmark generation must avoid paid APIs, preserve data governance, and remain repeatable as graphs evolve.

1 Introduction

Property graphs encode enterprise relationships directly: account transfers, identity entitlements, access paths, customer interactions, and fraud rings. Cypher is a common query language for these graphs. As LLMs become natural-language interfaces for graph analytics, organizations need reliable benchmarks for natural-language-to-Cypher systems on their own schemas.

Static public benchmarks are insufficient for this setting. They cannot contain private labels, relationship types, categorical values, or domain-specific wording, and they do not track schema evolution. PIPE-Cypher addresses this gap by generating private, repeatable, execution-grounded NL-Cypher benchmarks.

2 Related Work

Recent Text2Cypher resources, including Mind the Query (Chauhan et al., 2025), SyntheT2C (Zhong et al., 2025), Auto-Cypher (Tiwari et al., 2025), and Text2Cypher (Ozsoy et al., 2025), show that

high-quality Cypher data requires more than asking an LLM for query pairs. These works motivate schema grounding, execution checks, synthetic generation, verification, and complexity-aware evaluation. PIPE-Cypher differs by targeting private enterprise benchmark generation as the primary artifact: organizations bring their own graph, run local models, enforce Cypher constraints, and receive an executable benchmark with diversity and judge metadata.

Text-to-SQL benchmarks such as Spider 2.0 (Lei et al., 2024) and BIRD (Li et al., 2023) motivate realistic database tasks, execution-based evaluation, and enterprise workflow complexity. CIKM AutoQuery (Zheng et al., 2024) motivates strategy-level workload generation analysis and separating workload quality from model quality. LDBC FinBench (Qi et al., 2023) and SNB (Püroja et al., 2023) provide graph workloads suitable for industry-oriented evaluation.

For generated benchmark quality, PIPE-Cypher combines text-generation diversity diagnostics with text-to-query structure metrics. Distinct-n (Li et al., 2016) and self-BLEU-style redundancy checks (Zhu et al., 2018) capture lexical repetition, while schema coverage, query-signature diversity, structural feature rates, and normalized entropy over graph/category/difficulty cells capture properties that matter for executable Cypher benchmarks.

3 Method

PIPE-Cypher has six stages: schema profiling, template generation, reverse Cypher grounding, constrained Cypher generation and repair, deterministic validation and execution, and LLM-judge review. Deterministic checks enforce read-only safety, syntax shape, schema-visible labels and properties, explicit relationship types, observed relationship directions, schema-provided categorical values, and non-empty execution where required. A

lightweight Cypher analyzer extracts return aliases, variables, labels, relationship observations, risky constructs, and rewrite skip reasons so normalization is auditable rather than a silent string edit. The direction gate interprets both outgoing and incoming Cypher arrow syntax before schema checking, and rejects undirected relationship patterns so directionality is an executable constraint rather than only a prompt instruction.

4 Implementation

The implementation is a Python package with Neo4j as the experimental backend. The paper frames the contribution around Cypher and property graphs rather than a single vendor. Local models are served through a vLLM/OpenAI-compatible endpoint on ds-serv6, using Qwen3.5-9B for generation and judging.

All reported generation, judging, and downstream evaluation runs use a local Qwen3.5-9B endpoint. This keeps the study aligned with an enterprise deployment pattern in which the benchmark pipeline runs inside the organization’s compute boundary without paid generation APIs.

The built-in FinBench profile is grounded in the public datagen snapshot export and includes typed node and relationship properties. Live schema inspection also records low-cardinality string values as categorical constraints for prompting and validation. The generated Neo4j import script uses uniqueness constraints for nodes and creates, rather than merges, transaction relationships so repeated account-to-account events are preserved. The SNB reference profile is grounded in the official Neo4j/Cypher headers and read-query files.

For generation, PIPE-Cypher uses seeded workload templates with deterministic reverse-binding queries and a mixed mode that prepends proven workload seeds to LLM-proposed templates. Every run records accepted and rejected candidates for yield analysis. Retrieved few-shot examples replace graph-specific values with typed placeholders, preserving query structure while reducing tenant-value leakage and memorized entity reuse. A dependency-light value grounder adds typed prompt annotations for categorical values and reverse-bound entities, covering punctuation variants, possessives, plurals, synonyms, name partials, and small typos.

Inspired by Mind the Query’s prompt-setting analysis, the implementation exposes prompt pro-

files for schema-only, instructions-only, examples-only, examples-plus-instructions, and full governed PIPE-Cypher generation. These profiles are configured as ablation variants and must pass the same target-size and paper-readiness standards before any result is reported.

For LLM-judge review, the implementation slices schema context to the labels, relationship types, and properties used by the candidate query. This preserves validation against the full schema while keeping local 9B prompts within context limits on larger schemas.

The deterministic Cypher layer borrows from production lessons in the BalkanID Cypher system: schema-only prompting, exact matching, relationship direction discipline, RETURN DISTINCT, reserved variable rejection, categorical values, required contextual return columns, fuzzy value annotations, placeholderized retrieval examples, and parser-aware rewrite boundaries. PIPE-Cypher records parser-style structure features and skips rewrites for risky constructs such as UNION, CALL, UNWIND, WHERE EXISTS, multiple WHERE clauses, or reserved variables.

We also estimate seed-template capacity before full generation. This check caught and fixed a scale blocker: reverse-binding execution was hard-capped to 10 rows, and no-slot negation/ranking seeds could not support full category targets. PIPE-Cypher now uses the configured binding limit and includes slotted negation and ranking seeds for both graphs.

For arbitrary enterprise onboarding, PIPE-Cypher includes a deployment template for a company’s own graph, read-only credentials, local model endpoint, schema introspection, privacy policy, dry run, scaled run, audit, and redacted export. Configurable value-sampling policies bound which low-cardinality graph values enter schema prompts, and redacted exports replace quoted literals, entity values, and string-valued result samples with stable placeholders for broader internal review.

5 Experiments

The generated benchmark contains 3,000 accepted examples: 2,000 from LDBC FinBench and 1,000 from LDBC SNB. Categories include simple retrieval, complex retrieval, simple aggregation, complex aggregation, boolean existence, negation/difference, path/temporal transaction, and ranking/top-k. Appendix Table 8 reports graph

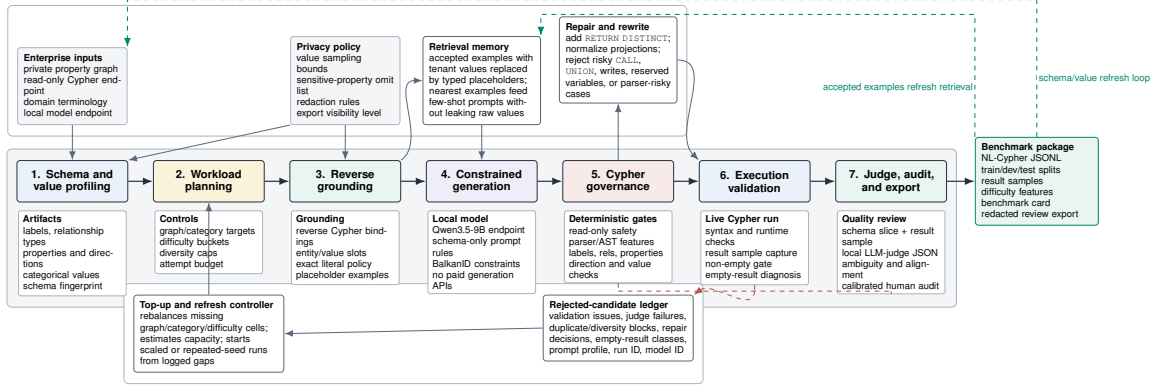


Figure 1: PIPE-Cypher benchmark generation pipeline. The key industry additions beyond static Text2Cypher dataset construction are privacy/value policies, reverse grounding, Cypher governance, execution diagnostics, local judge calibration, and benchmark export/refresh.

Gate	Evidence checked	Failure mode caught
Read-only safety	blocked Cypher write/admin tokens	destructive or operational queries
Schema validity	labels, relationship types, properties, categorical values	hallucinated graph vocabulary or values
Direction validity	observed relationship directions	reversed or invalid traversals
Cypher analysis and normalization	structure extraction, rewrite safety, RETURN DISTINCT	unsafe rewrites or duplicate/noisy rows
Question constraints	quoted values use exact literals	fuzzy matching of exact user values
Execution	query runs and returns rows	syntactic validity without answerability
LLM judge	ambiguity, semantic alignment, schema use	executable but semantically weak examples

Table 1: PIPE-Cypher candidate acceptance gates.

Graph	Target/category	Binding limit	Min seed capacity	All categories meet target
FinBench	250	300	300	Yes
SNB	125	200	200	Yes

Table 2: Built-in seed-template capacity after removing the reverse-binding 10-row cap and adding slotted negation/ranking seeds. Capacity is a theoretical scale-readiness check; final examples still must pass validation, execution, diversity, and judge gates.

Setting	Value
Accepted examples	3,000
Primary graph	LDBC FinBench, 2,000 examples
Secondary graph	LDBC SNB, 1,000 examples
Categories	8 balanced categories, 375 each
Generation/judge model	Local Qwen3.5-9B
Execution backend	Neo4j Community, two live databases
Judge audit packet	80 sampled accepted/rejected pairs

Table 3: Full live experimental setup used for the generated benchmark artifact.

Graph	Candidates	Accepted	Acceptance	Categories at target
FinBench	3,404	2,000	0.588	8/8
SNB	1,373	1,000	0.728	8/8
Total	4,777	3,000	0.628	16/16

Table 4: Full live generation with local Qwen3.5-9B. Candidate counts include the initial sequential run and category-specific recovery top-ups.

statistics and the current ICIJ onboarding status.

Downstream Text2Cypher evaluation uses execution accuracy and answer-set precision/recall/F1 as primary correctness metrics. The benchmark evaluator also supports supplementary reference-based text metrics over serialized answer sets and query strings, including ROUGE, BLEU, METEOR, BERTScore, FrugalScore, cosine similarity, Jaro-Winkler similarity, and exact match; Appendix Table 20 lists the supported metrics and citations. We treat these as debugging and near-match diagnostics rather than substitutes for executable correctness.

6 Results

The full live run produced 3,000 accepted examples from 4,777 candidates using local Qwen3.5-9B for

generation and judging. Category-specific recovery top-ups filled the only under-target categories from the initial sequential run. Every exported example passed read-only, syntax, schema, execution, non-empty result, and judge gates.

The accepted full-run records were exported into a benchmark package with stable example identifiers, train/dev/test splits, result samples, gate meta-

Export artifact	Examples	FinBench	SNB	Train	Dev	Test
Live full benchmark	3,000	2,000	1,000	2,408	296	296

Table 5: Accepted live full benchmark package with stable IDs, gate metadata, result samples, statistics, and manifest hash 8bc79a53a06b291a.

Artifact property	Value
Categories	8 balanced categories
FinBench/category	250
SNB/category	125
Difficulty split	1,521 easy / 1,479 medium
Unique labels / rel. types	7 / 9
Read/syntax/schema/exec/judge gates	3,000/3,000

Table 6: Distribution and gate summary for the exported full benchmark artifact.

data, aggregate statistics, and a manifest hash for artifact tracking. The export contains 3,000 accepted examples: 2,000 FinBench, 1,000 SNB, and 375 accepted examples in every planned category.

We report diversity as a diagnostic, not only as a design target. The full export has perfect category balance and near-perfect difficulty balance, uses 1,115 unique grounded entity values, and exactly quotes grounded values in 82.6% of examples with entity bindings. The reported local-model run still has low query-signature diversity because seeded graph-grounded templates are doing substantial work. This is useful evidence for industry deployment: the artifact is balanced and executable, while the appendix makes template and value concentration visible rather than hiding it.

The full-run failure taxonomy is concentrated in diversity/duplicate controls during recovery and empty execution results; only 2/4,777 candidates were schema-invalid after deterministic Cypher governance.

For judge calibration, the released artifacts include an 80-row audit packet sampled from accepted and rejected full-run candidates, with a 40/40 judge accept/reject split and coverage over both graphs and all eight categories. A completed human audit shows 80.0% agreement, Cohen’s $\kappa = 0.60$, judge precision and specificity of 1.00, recall of 0.714, and no false accepts in the sample. The judge is therefore conservative: it protects accepted-example quality while rejecting some human-approved candidates and reducing yield. Human labels are used only for post-hoc calibration, not as a generation gate.

Finally, we ran a downstream Text2Cypher evaluation using local Qwen3.5-9B on the exported

Split	Examples	Parse	Schema	Exec. success	Exec. acc.	Answer F1
Live full test	296	0.959	0.905	0.622	0.189	0.189

Table 7: Downstream Text2Cypher evaluation for local Qwen3.5-9B on the exported full benchmark test split.

full test split. The model saw schema text and the natural-language question, generated Cypher, and was evaluated by live execution against the corresponding FinBench or SNB database. The result verifies that the exported artifact can drive downstream model evaluation and gives a difficulty signal: the zero-shot 9B baseline solves simple aggregation examples but struggles on more operational categories.

7 Industry Use

PIPE-Cypher is intended to be rerun when an enterprise graph changes. The generated artifact records the schema snapshot, graph profile, model identifier, validation gates, execution samples, judge scores, difficulty features, and source run for every example. Long-running jobs can be monitored from JSONL records and recovered with category-specific top-up runs that reject questions accepted in earlier runs. This design supports private benchmark refreshes without exposing schemas or values to paid APIs, while still leaving audit hooks for data owners to sample accepted and rejected examples.

The deployment path is designed for arbitrary enterprise schemas: teams provide read-only graph credentials, introspect labels/relationships/properties, choose a bounded categorical-value sampling policy, connect a local model endpoint, run a dry pass, scale generation, calibrate the judge, and export either raw internal artifacts or redacted review artifacts. This makes the FinBench/SNB study a reproducible evaluation setting rather than a hard-coded graph assumption.

8 Conclusion

PIPE-Cypher provides a practical path for enterprises to create private, executable, balanced NL-to-Cypher benchmarks. The pipeline combines schema introspection, constrained generation, deterministic validation, execution feedback, diversity control, and automated judge review.

Limitations

Execution validity does not guarantee semantic correctness. The completed audit suggests a conservative judge with no false accepts in the labeled

sample, but larger audits may reveal additional failure modes. Generated artifacts may contain private schema details or sampled values, so organizations should apply privacy redaction and normal data governance controls before broad sharing.

Ethics Statement

PIPE-Cypher is designed for private benchmark generation under local-model deployment constraints. Organizations using it should restrict benchmark artifacts to authorized users, review sampled values for sensitive content, and document whether judge calibration relied on human annotators.

References

Satanjeev Banerjee and Alon Lavie. 2005. [METEOR: An automatic metric for MT evaluation with improved correlation with human judgments](#). In *Proceedings of the ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization*, pages 65–72, Ann Arbor, Michigan. Association for Computational Linguistics.

Vashu Chauhan, Shobhit Raj, Shashank Mujumdar, Avirup Saha, and Anannay Jain. 2025. [Mind the query: A benchmark dataset towards Text2Cypher task](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: Industry Track*, pages 1890–1905, Suzhou, China. Association for Computational Linguistics.

Matthew A. Jaro. 1989. [Advances in record-linkage methodology as applied to matching the 1985 census of tampa, florida](#). *Journal of the American Statistical Association*, 84(406):414–420.

Moussa Kamal Eddine, Guokan Shang, Antoine Tixier, and Michalis Vazirgiannis. 2022. [FrugalScore: Learning cheaper, lighter and faster evaluation metrics for automatic text generation](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1305–1318, Dublin, Ireland. Association for Computational Linguistics.

Fangyu Lei, Jixuan Chen, Yuxiao Ye, Ruisheng Cao, Dongchan Shin, Hongjin Su, Zhaoqing Suo, Hongcheng Gao, Wenjing Hu, Pengcheng Yin, Victor Zhong, Caiming Xiong, Ruoxi Sun, Qian Liu, Sida Wang, and Tao Yu. 2024. [Spider 2.0: Evaluating language models on real-world enterprise text-to-SQL workflows](#). *Preprint*, arXiv:2411.07763.

Jinyang Li, Binyuan Hui, Ge Qu, Jiayi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Rongyu Cao, Ruiying Geng, Nan Huo, Xuanhe Zhou, Chenhao Ma, Guoliang Li, Kevin C. C. Chang, Fei Huang,

Reynold Cheng, and Yongbin Li. 2023. [Can LLM already serve as a database interface? a BIG bench for large-scale database grounded text-to-SQLs](#). In *Advances in Neural Information Processing Systems*.

Jiwei Li, Michel Galley, Chris Brockett, Jianfeng Gao, and Bill Dolan. 2016. [A diversity-promoting objective function for neural conversation models](#). In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 110–119, San Diego, California. Association for Computational Linguistics.

Chin-Yew Lin. 2004. [ROUGE: A package for automatic evaluation of summaries](#). In *Text Summarization Branches Out*, pages 74–81, Barcelona, Spain. Association for Computational Linguistics.

Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. 2008. *Introduction to Information Retrieval*. Cambridge University Press.

Makbule Gulcin Ozsoy, Leila Messallem, Jon Besga, and Gianandrea Minneci. 2025. [Text2Cypher: Bridging natural language and graph databases](#). In *Proceedings of the Workshop on Generative AI and Knowledge Graphs*, pages 100–108, Abu Dhabi, UAE. International Committee on Computational Linguistics.

Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. [Bleu: a method for automatic evaluation of machine translation](#). In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318, Philadelphia, Pennsylvania, USA. Association for Computational Linguistics.

David Püroja, Jack Waudby, Peter Boncz, and Gábor Székely. 2023. [The LDBC social network benchmark interactive workload v2: A transactional graph query benchmark with deep delete operations](#). *Preprint*, arXiv:2307.04820.

Shipeng Qi, Heng Lin, Zhihui Guo, Gábor Székely, Bing Tong, Yan Zhou, Bin Yang, Jiansong Zhang, Zheng Wang, Youren Shen, Changyuan Wang, Parviz Peiravi, Henry Gabb, and Ben Steer. 2023. [The LDBC financial benchmark](#). *Preprint*, arXiv:2306.15975.

Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. 2016. [SQuAD: 100,000+ questions for machine comprehension of text](#). In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392, Austin, Texas. Association for Computational Linguistics.

Aman Tiwari, Shiva Krishna Reddy Malay, Vikas Yadav, Masoud Hashemi, and Sathwik Tejaswi Madhusudhan. 2025. [Auto-cypher: Improving LLMs on](#)

Graph	Nodes	Relationships	Labels	Rel. types	Paper status
FinBench	10,006	57,622	5	9	reported
SNB	34,735	70,842	14	15	reported
ICIJ Offshore Leaks	–	–	5	14	onboarding only

Table 8: Study graph statistics. ICIJ is listed as an onboarding workload only until its live generation run passes the same audit standard as FinBench and SNB.

cypher generation via LLM-supervised generation-verification framework. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2: Short Papers*, pages 623–640, Albuquerque, New Mexico. Association for Computational Linguistics.

William E. Winkler. 1990. [String comparator metrics and enhanced decision rules in the Fellegi-Sunter model of record linkage](#). In *Proceedings of the Section on Survey Research Methods*, pages 354–359. American Statistical Association.

Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. 2020. [BERTScore: Evaluating text generation with BERT](#). In *International Conference on Learning Representations*.

Xiuwen Zheng, Arun Kumar, and Amarnath Gupta. 2024. [Generating cross-model analytics workloads using LLMs](#). In *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management*, pages 4303–4307. ACM.

Zijie Zhong, Linqing Zhong, Zhaoze Sun, Qingyun Jin, Zengchang Qin, and Xiaofan Zhang. 2025. [SyntheT2C: Generating synthetic data for fine-tuning large language models on the Text2Cypher task](#). In *Proceedings of the 31st International Conference on Computational Linguistics*, pages 672–692, Abu Dhabi, UAE. Association for Computational Linguistics.

Yaoming Zhu, Sidi Lu, Lei Zheng, Jiaxian Guo, Weinan Zhang, Jun Wang, and Yong Yu. 2018. [Texygen: A benchmarking platform for text generation models](#). In *The 41st International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 1097–1100. ACM.

A Additional Results

B Claim–Evidence Map

This section maps the main paper claims to the strongest current evidence and to the remaining risks. This is intended as a reviewer-facing audit surface: claims with pending evidence remain marked as such rather than being folded into the main results.

1. **Claim 1.** PIPE-Cypher can generate a private, executable NL-to-Cypher benchmark

Metric	Value
Question Distinct-1	0.041
Question Distinct-2	0.088
Question self-BLEU-2 (sampled)	0.826
Unique query-signature ratio	0.010
Top query-signature share	0.083
Category normalized entropy	1.000
Graph-category normalized entropy	0.980
Difficulty normalized entropy	1.000
Label coverage	0.412
Relationship-type coverage	0.375
Property-name coverage	0.407
Unique grounded-value ratio	0.373
Grounded values exactly quoted	0.826
Aggregation / negation / ordering rates	0.500 / 0.125 / 0.125

Table 9: Diversity diagnostics for the full exported benchmark. Distinct-n follows text-generation usage; self-BLEU is lower when questions are less redundant; grounded-value metrics are aggregate-only and do not list raw values.

Failure bucket	Count	Share of rejected
Diversity/duplicate control	1,335	0.751
Empty result	374	0.210
Judge semantic reject	66	0.037
Schema invalid	2	0.001

Table 10: Failure taxonomy over full-run generation candidates before benchmark export. Accepted examples are excluded from the bucket shares.

over enterprise-style property graphs using local models.

Evidence. The full local Qwen3.5-9B run exported 3,000 accepted examples: 2,000 FinBench and 1,000 SNB, with 375 examples in every planned workload category and all accepted examples passing read-only, syntax, schema, execution, non-empty-result, and judge gates.

Key artifacts. benchmark export; manifest snapshot; paper table: benchmark export; paper table: full results

Status. Supported by live local-model evidence.

Risk. Generation quality may change with private enterprise schemas beyond FinBench/SNB.

2. **Claim 2.** Cypher-specific governance makes generated benchmark candidates safer and more auditable than prompt-only generation.

Evidence. The pipeline validates read-only safety, schema-visible labels, node and relationship properties, relationships, observed

PIPE-Cypher category	Closest MTQ category	Added enterprise distinction
Simple retrieval	SR	Direct node/edge lookup with exact filters.
Complex retrieval	CR	Multi-hop or multi-pattern retrieval.
Simple aggregation	SA	Single aggregation such as count/min/max/average.
Complex aggregation	CA	Grouped or multi-stage aggregation over graph neighborhoods.
Boolean existence	EQ	Precise yes/no or existence answer.
Negation/difference	CR	Absence, anti-join, or difference query.
Path/temporal transaction	CR/CA	Temporal or path-oriented transaction neighborhood.
Ranking/top-k	SA/CA	Ordered top-k query with explicit limit.

Table 11: Category crosswalk to Mind the Query. PIPE-Cypher keeps the familiar retrieval/aggregation/evaluation-query structure while adding enterprise workloads such as negation, temporal paths, and ranking.

Gate or ledger	Count	Denominator
Export accepted	3,000	3,000
Read-only safety	3,000	3,000
Syntax validity	3,000	3,000
Schema/value validity	3,000	3,000
Execution success	3,000	3,000
LLM judge pass	3,000	3,000
Rejected candidates logged	1,777	4,777

Table 12: PIPE-Cypher validation cascade for the full export and its logged candidate ledger. Unlike Mind the Query, human review is calibration-only.

relationship direction, categorical values, contextual return columns, parser-style structure, execution success, and LLM-judge review. In the full run, only 2 of 4,777 candidates were schema-invalid after deterministic governance, and all 3,000 exported examples passed the deterministic and judge gates.

Key artifacts. code: validator.py; code: cypher_parser.py; code: question_constraints.py; snapshot: failure taxonomy; +1 more

Status. Supported by code, tests, and full-run failure taxonomy.

Risk. Semantic correctness still depends on execution evidence and judge review; the completed 80-row audit is a calibration sample rather than an exhaustive proof.

- Claim 3.** Reverse grounding and structured workload seeds are designed to improve reliability under local-model constraints.

Evidence. The target-50 FinBench/SNB ablation suite completed 14/14 graph/variant cells, reached every category target for all non-empty/non-unconstrained runs, and passed the paper-readiness audit with logs, model IDs, code revision, run summaries, and gate-rate availability. Appendix-only tables plus yield and gate-rate figures report the au-

Profile	Added constraint	Intended evidence
Schema-only	Use only visible schema.	Baseline for schema-grounding prompting.
Instructions	Exact values, datatype rules, no nested aggregations.	Targets MTQ-style prompt refinements.
Examples	Placeholderized retrieval NL-Cypher pairs.	Tests whether examples help without extra governance.
Examples + instructions	Few-shot plus explicit rules.	Closest controlled analogue to MTQ Table 5.
Full governed	BalkanID-derived hints, AST-safe rewrites, judge gate.	PIPE-Cypher production setting.

Table 13: Prompt-profile plan inspired by Mind the Query’s prompt-variant analysis. Only completed, audited target-50-or-larger results should be promoted into results tables.

Dimension	Mind the Query	PIPE-Cypher
Generation review gate	Manual logical review	Deterministic gates + local LLM judge
Human effort	Reported 1,400 person-hours	80-row post-hoc judge calibration audit
Private values	Public benchmark values	Configurable sampling and redacted export
Refresh	Static dataset release	ReRunnable private benchmark factory
Model endpoint	Gemini for generation	Local Qwen3.5-9B endpoint

Table 14: Industry deployment contrast with Mind the Query. PIPE-Cypher focuses on private refreshable benchmark generation rather than a one-time public dataset.

dated target-50 evidence while the target-100 follow-up and seeded target-50 repeat continue for stronger final reliability and variance claims.

Key artifacts. script: run live ablation suite; script: monitor remote ablation suite; script: monitor remote ablation queue; script: collect remote ablation suite; +9 more

Status. Supported by audited target-50 appendix evidence; target-100 and seeded target-50 repeat runs are still pending for stronger final reliability claims.

Risk. A single target-50 suite is the minimum publishable ablation scale, not the strongest possible evidence. If target-100 or repeated-seed runs complete, final claims should use those sensitivity results instead of relying only on target-50.

- Claim 4.** The benchmark controls category and difficulty balance while exposing residual diversity limits.

Evidence. The full export has perfect category balance, planned 2:1 FinBench/SNB graph mix, and a near-even easy/medium split. Diversity diagnostics report Distinct-n, sampled self-BLEU-2, query-signature diversity, normalized entropy, schema coverage, structural feature rates, and aggregate value-grounding diagnostics. The full export uses 1,115 unique grounded entity values and exactly quotes grounded values in 82.6% of examples with entity bindings, making template and value concentration visible instead of hidden.

Setting	Graph	Records	Accepted	Acceptance	Categories at target
Unconstrained LLM	FinBench	0	0	0.000	0/8
Unconstrained LLM	SNB	0	0	0.000	0/8
Reverse-only	FinBench	400	400	1.000	8/8
Reverse-only	SNB	402	400	0.995	8/8
Validators+repair	FinBench	400	400	1.000	8/8
Validators+repair	SNB	402	400	0.995	8/8
No retrieval	FinBench	429	400	0.932	8/8
No retrieval	SNB	402	400	0.995	8/8
No rewrite	FinBench	423	400	0.946	8/8
No rewrite	SNB	402	400	0.995	8/8
No LLM judge	FinBench	401	400	0.998	8/8
No LLM judge	SNB	406	400	0.985	8/8
Full PIPE-Cypher	FinBench	416	400	0.962	8/8
Full PIPE-Cypher	SNB	402	400	0.995	8/8

Table 15: Live target-50 ablation evidence with local Qwen3.5-9B. Each graph run targets 50 accepted examples per category.

Setting	Graph	Read-only	Syntax	Schema	Exec.	Judge
Unconstrained LLM	FinBench	0.000	0.000	0.000	0.000	0.000
Unconstrained LLM	SNB	0.000	0.000	0.000	0.000	0.000
Reverse-only	FinBench	1.000	1.000	1.000	1.000	1.000
Reverse-only	SNB	1.000	1.000	0.995	0.995	0.995
Validators+repair	FinBench	1.000	1.000	1.000	1.000	1.000
Validators+repair	SNB	1.000	1.000	0.995	0.995	0.995
No retrieval	FinBench	1.000	1.000	1.000	1.000	0.932
No retrieval	SNB	1.000	1.000	0.995	0.995	0.995
No rewrite	FinBench	1.000	1.000	1.000	1.000	0.946
No rewrite	SNB	1.000	1.000	0.995	0.995	0.995
No LLM judge	FinBench	1.000	1.000	1.000	1.000	0.998
No LLM judge	SNB	1.000	1.000	0.995	0.995	0.985
Full PIPE-Cypher	FinBench	1.000	1.000	1.000	1.000	0.962
Full PIPE-Cypher	SNB	1.000	1.000	0.995	0.995	0.995

Table 16: Quality-gate rates for the live target-50 ablation suite. Rates are computed over all generated records in each graph/setting.

Audit packet property	Value
Rows	80
Judge accept / reject	40 / 40
FinBench / SNB rows	48 / 32
Easy / medium rows	26 / 54
Labeled rows	80
Agreement / Cohen’s κ	0.800 / 0.600
Judge precision / recall	1.000 / 0.714
Judge specificity	1.000
False accepts / false rejects	0 / 16

Table 17: Post-hoc judge calibration from the completed 80-row human audit. Labels are used only to calibrate the automated judge, not as a generation gate.

Metric	Point	95% CI	SE	N
Parse valid	0.959	[0.936, 0.980]	0.012	296
Schema valid	0.905	[0.872, 0.939]	0.017	296
Execution success	0.622	[0.564, 0.676]	0.028	296
Execution accuracy	0.189	[0.145, 0.233]	0.022	296
Answer F1	0.189	[0.149, 0.236]	0.023	296

Table 18: Bootstrap uncertainty for downstream Text2Cypher evaluation on the full exported test split. Intervals use row-level resampling with a fixed seed and are intended for appendix reporting.

Key artifacts. snapshot: diversity report; paper table: diversity; paper figure: diversity diagnostics; paper figure: full export distribution

Status. Supported by full-export statistics and diversity diagnostics.

Risk. Low query-signature diversity in the reported local-model run remains a limitation and ablation target.

- Claim 5.** The generated benchmark is useful for downstream Text2Cypher evaluation.

Evidence. A zero-shot local Qwen3.5-9B downstream evaluation over the 296-example full test split reports parse validity, schema validity, execution success, execution accuracy, answer F1, per-category breakdowns, fixed-seed bootstrap uncertainty intervals, and a row-level error taxonomy. The model reaches 0.189 execution accuracy with a 95% bootstrap interval of [0.145, 0.233] and 0.622 execution success with interval [0.564, 0.676]. The error taxonomy classifies 240 incorrect rows into answer mismatches, execution failures, schema invalid predictions, and parse-invalid predictions, showing the benchmark can discriminate easy aggregation from harder operational categories and can explain model

failure modes.

Key artifacts. downstream summary; snapshot: downstream uncertainty; snapshot: downstream error report; research note: downstream evaluation; +6 more

Status. Supported by live downstream evaluation against FinBench/SNB databases, with bootstrap uncertainty computed from row-level outputs.

Risk. Only one downstream model has been evaluated so far.

- Claim 6.** Automated LLM-judge review can replace human-in-the-loop generation gates while still exposing judge risk.

Evidence. The pipeline uses deterministic validation plus LLM-judge review as the generation gate. A post-hoc 80-row judge audit packet preserves 40 judge accepts, 40 judge rejects, both graphs, and all eight categories. The completed human audit reports 80.0% agreement, Cohen’s kappa 0.60, balanced accuracy 0.857, judge precision 1.00, judge specificity 1.00, judge recall 0.714, zero false accepts, and 16 false rejects. The analyzer

Downstream outcome	Count	Share of incorrect
Answer mismatch	128	0.533
Execution failed	72	0.300
Schema invalid	28	0.117
Parse invalid	12	0.050

Table 19: Downstream Text2Cypher failure taxonomy for local Qwen3.5-9B on the full exported test split. Shares exclude exact-answer matches.

requires complete labels for paper-facing calibration, and the tracked summary avoids committing raw value-bearing audit rows.

Key artifacts. code: judge.py; code: calibration.py; code: judge_audit_packet.py; script: create judge annotation sheets; +6 more

Status. Supported by completed human audit summary.

Risk. The judge appears conservative on this sample; larger audits or different enterprise schemas may reveal false accepts.

- Claim 7.** PIPE-Cypher supports arbitrary enterprise schema onboarding with configurable privacy redaction and value-sampling policies.

Evidence. The repo includes an enterprise deployment guide, an enterprise template config, privacy config fields, schema introspection that reads categorical-value thresholds, maximum sampled value length, and omitted-property patterns from config, deterministic redaction helpers, and a redacted benchmark export CLI. The redaction policy replaces quoted Cypher literals, question mentions, entity values, and string-valued result samples with stable placeholders while preserving query structure and gate metadata. The repo now also includes an ICIJ Offshore Leaks onboarding profile, config, download/load helpers, and graph plan for a public finance/compliance property graph beyond the LDBC study workloads.

Key artifacts. research note: enterprise deployment guide; research note: icij offshore-leaks graph plan; enterprise_template.yaml; icij_offshoreleaks_smoke.yaml; +10 more

Status. Supported by docs, config, code, deterministic tests, and a concrete third-graph onboarding plan; live ICIJ generation remains pending.

Risk. ICIJ is public rather than private/anonymized enterprise data. The strongest industry validation would still come from a real tenant graph or a completed audited ICIJ live run.

C Prompt Contracts

PIPE-Cypher treats prompts as versioned implementation artifacts. The list summarizes the prompt contracts used for generation, repair, judging, and downstream evaluation; hashes fingerprint the full prompt constants in the codebase.

- Template generation.** Stage: Workload proposal. SHA-256: 10b25b.

Contract. Schema-only labels, relationships, properties, and categorical values; Realistic enterprise analyst wording; At most two typed slots and JSON-only output

- Reverse binding.** Stage: Graph grounding. SHA-256: 48e949.

Contract. Read-only MATCH/WHERE/RETURN DISTINCT/LIMIT only; Slot variables named exactly as requested; Forward relationship directions from the schema

- Cypher generation.** Stage: Candidate query. SHA-256: 61c557.

Contract. Only schema-visible constructs and observed directions; RETURN DISTINCT for set returns and exact equality for quoted values; Context columns, categorical hints, placeholderized retrieval, and no writes

- Repair.** Stage: Validation feedback. SHA-256: 56c419.

Contract. Preserve question intent while fixing validation or execution issues; Keep query read-only and schema-grounded; Return only corrected Cypher

- LLM judge.** Stage: Quality gate. SHA-256: 421c7b.

Contract. Inputs include question, Cypher, schema slice, execution rows, and validation summary; Strict JSON scores for ambiguity, semantic alignment, schema use, and difficulty; Categorical values constrain query literals, not observed result-row values; Pass only

Metric	PIPE-Cypher support	Primary citation
ROUGE-1/2/L	Reference overlap over serialized answer sets and query strings	(Lin, 2004)
BLEU	Sentence-level reference overlap with brevity penalty	(Papineni et al., 2002)
METEOR	Exact-token METEOR-style precision/recall with fragmentation penalty	(Banerjee and Lavie, 2005)
BERTScore	Optional contextual-embedding similarity when installed	(Zhang et al., 2020)
FrugalScore	Optional efficient learned NLG metric when installed	(Kamal Eddine et al., 2022)
Cosine	Token-count vector cosine similarity	(Manning et al., 2008)
Jaro-Winkler	Character-level string similarity for near-match diagnostics	(Jaro, 1989; Winkler, 1990)
Exact match	Raw and normalized string exact match	(Rajpurkar et al., 2016)

Table 20: Supplementary reference-based text metrics supported by the PIPE-Cypher evaluator. These metrics are useful for debugging answer rendering, paraphrase sensitivity, and near-match behavior, but they do not replace execution accuracy or answer-set F1 for Text2Cypher correctness. BERTScore and FrugalScore are optional integrations because they require additional metric/model packages.

useful, unambiguous enterprise benchmark examples

```
(:Account)
RETURN DISTINCT SUM(t.amount) AS
TotalTransferredAmount
```

6. **Downstream Text2Cypher.** Stage: Model evaluation. SHA-256: 4c07ff.
- Contract.* Read-only Cypher only; Schema-visible constructs and exact direction preservation; RETURN DISTINCT, count/ranking rules, and no explanations

Structure: join_heavy, aggregation.

Relationships: OWN_ACCOUNT, TRANSFER_TO.

Gates: RO/Syn/Schema/Exec/Judge.

Result sample: {TotalTransferredAmount: 14972318.41}; observed rows: 1

D Representative Accepted Examples

The examples below are selected in stable identifier order from the tracked full-export snapshot, one per graph/category cell when available. They show the natural-language question, accepted Cypher, structural tags, gate status, and a bounded execution-result sample.

1. **FinBench / Boolean Existence / easy.** *Question:* Does account '187743809466009406' have any outgoing transfer?

```
MATCH (src:Account {accountId:
'187743809466009406'})
-[:TRANSFER_TO]->
(:Account)
RETURN DISTINCT COUNT(DISTINCT src)
> 0
AS HasOutgoingTransfer
```

Structure: single_hop, aggregation.

Relationships: TRANSFER_TO.

Gates: RO/Syn/Schema/Exec/Judge.

Result sample: {HasOutgoingTransfer: True}; observed rows: 1

2. **FinBench / Complex Aggregation / medium.** *Question:* What is the total transferred amount from accounts owned by person 'Gwar'?

```
MATCH (p:Person {personName:
'Gwar'})
-[:OWN_ACCOUNT]->
(src:Account)-[t:TRANSFER_TO]->
```

3. **FinBench / Complex Retrieval / easy.** *Question:* Which accounts received transfers from accounts owned by person 'Zof'?

```
MATCH (p:Person {personName:
'Zof'})
-[:OWN_ACCOUNT]-> (src:Account)
-[:TRANSFER_TO]-> (dst:Account)
RETURN DISTINCT dst.accountId AS
AccountId, dst.accountType AS
AccountType, dst.isBlocked AS
IsBlocked
LIMIT 300
```

Structure: join_heavy, bounded_result.

Relationships: OWN_ACCOUNT, TRANSFER_TO.

Gates: RO/Syn/Schema/Exec/Judge.

Result sample: {AccountId: 4687402787162554854, AccountType: credit card, IsBlocked: False}; observed rows: 3

4. **FinBench / Negation Difference / medium.** *Question:* Which accounts owned by person 'Kant' have not sent any transfers?

```
MATCH (p:Person {personName:
'Kant'})
-[:OWN_ACCOUNT]-> (a:Account)
WHERE NOT (a) -[:TRANSFER_TO]->
(:Account)
RETURN DISTINCT a.accountId AS
AccountId, a.accountType AS
AccountType,
a.isBlocked AS IsBlocked
LIMIT 300
```

733	<i>Structure:</i>	join_heavy, negation,	internet account, IsBlocked: False}; observed	788
734		bounded_result.	rows: 1	789
735	<i>Relationships:</i>	OWN_ACCOUNT, TRANS-		790
736		FER_TO.		791
737	<i>Gates:</i>	RO/Syn/Schema/Exec/Judge.		792
738	<i>Result sample:</i>	{AccountId:		793
739		4739757132830738341, AccountType:		794
740		merchant account, IsBlocked: False};		795
741		observed rows: 1		796
742				797
743	5. FinBench / Path Temporal / medium. Question:	Which accounts can receive money		
744		within two transfer hops from accounts owned		
745		by person 'Sossamon'?		
746		<pre>MATCH (p:Person {personName: 747 'Sossamon'}) -[:OWN_ACCOUNT]-> 748 (src:Account) -[:TRANSFER_TO*1..2]-> 749 (dst:Account) 750 RETURN DISTINCT dst.accountId AS 751 AccountId, dst.accountType AS 752 AccountType, dst.isBlocked AS 753 IsBlocked 754 LIMIT 300</pre>		
755	<i>Structure:</i>	join_heavy, path, bounded_result.		
756	<i>Relationships:</i>	OWN_ACCOUNT, TRANS-		
757		FER_TO.		
758	<i>Gates:</i>	RO/Syn/Schema/Exec/Judge.		
759	<i>Result sample:</i>	{AccountId:		
760		4687402787162554854, AccountType:		
761		credit card, IsBlocked: False}; observed rows:		
762		4		
763	6. FinBench / Ranking Topk / medium. Question:	For accounts owned by person 'Barry',		
764		which account sent the highest total transfer		
765		amount?		
766				
767		<pre>MATCH (p:Person {personName: 768 'Barry'}) 769 -[:OWN_ACCOUNT]-> 770 (src:Account) -[:TRANSFER_TO]-> 771 (:Account) 772 WITH src, SUM(t.amount) AS 773 totalAmount 774 RETURN DISTINCT src.accountId AS 775 AccountId, src.accountType AS 776 AccountType, src.isBlocked AS 777 IsBlocked, 778 totalAmount 779 ORDER BY totalAmount DESC 780 LIMIT 1</pre>		
781	<i>Structure:</i>	join_heavy, aggregation, or-		
782		der_rank, bounded_result.		
783	<i>Relationships:</i>	OWN_ACCOUNT, TRANS-		
784		FER_TO.		
785	<i>Gates:</i>	RO/Syn/Schema/Exec/Judge.		
786	<i>Result sample:</i>	{AccountId:		
787		4732438783436260772, AccountType:		
	7. FinBench / Simple Aggregation / easy. Question:	How many accounts are owned by per-		
		son 'Kaewsuktae'?		
		<pre>MATCH (p:Person {personName: 793 'Kaewsuktae'}) -[:OWN_ACCOUNT]-> 794 (a:Account) 795 RETURN DISTINCT COUNT(DISTINCT a) AS 796 AccountCount 797</pre>		
	<i>Structure:</i>	single_hop, aggregation.		798
	<i>Relationships:</i>	OWN_ACCOUNT.		799
	<i>Gates:</i>	RO/Syn/Schema/Exec/Judge.		800
	<i>Result sample:</i>	{AccountCount: 1}; observed		801
		rows: 1		802
	8. FinBench / Simple Retrieval / easy. Question:	Which accounts are owned by person		803
		'Barry'?		804
				805
		<pre>MATCH (p:Person {personName: 806 'Barry'}) 807 -[:OWN_ACCOUNT]-> (a:Account) 808 RETURN DISTINCT a.accountId AS 809 AccountId, a.accountType AS 810 AccountType, 811 a.isBlocked AS IsBlocked 812 LIMIT 300 813</pre>		
	<i>Structure:</i>	single_hop, bounded_result.		814
	<i>Relationships:</i>	OWN_ACCOUNT.		815
	<i>Gates:</i>	RO/Syn/Schema/Exec/Judge.		816
	<i>Result sample:</i>	{AccountId:		817
		4732438783436260772, AccountType:		818
		internet account, IsBlocked: False}; observed		819
		rows: 1		820
	9. SNB / Boolean Existence / medium. Question:	Does person with id 6597069766828		821
		like any post?		822
		<pre>MATCH (p:Person {id: 824 6597069766828}) 825 OPTIONAL MATCH (p) -[:LIKES]-> 826 (post:Post) 827 RETURN DISTINCT COUNT(post) > 0 AS 828 LikesAnyPost 829</pre>		
	<i>Structure:</i>	single_hop, aggregation, optional.		830
	<i>Relationships:</i>	LIKES.		831
	<i>Gates:</i>	RO/Syn/Schema/Exec/Judge.		832
	<i>Result sample:</i>	{LikesAnyPost: True}; ob-		833
		served rows: 1		834
	10. SNB / Complex Aggregation / medium. Question:	How many distinct posts		835
		are in forums joined by person with id		836
		6597069766845?		837
				838

	MATCH (forum:Forum) -[:HAS_MEMBER]-> (p:Person {id: 6597069766845}) MATCH (forum) -[:CONTAINER_OF]-> (post:Post) RETURN DISTINCT COUNT(DISTINCT post) AS JoinedForumPostCount	AS LastName LIMIT 200	
	Structure: join_heavy, aggregation.	Structure: single_hop, path, bounded_result.	
	Relationships: CONTAINER_OF, HAS_MEMBER.	Relationships: KNOWS.	
	Gates: RO/Syn/Schema/Exec/Judge.	Gates: RO/Syn/Schema/Exec/Judge.	
	Result sample: {JoinedForumPostCount: 1}; observed rows: 1	Result sample: {FirstName: Rafael, LastName: Fernández, PersonId: 4398046511333}; observed rows: 25	
11. SNB / Complex Retrieval / easy. Question:	Which people are members of forums contain- ing posts tagged 'Manuel_Noriega'?	14. SNB / Ranking Topk / medium. Ques- tion: Among forums tagged 'James_K._Polk', which forum has the most members?	
	MATCH (forum:Forum) -[:HAS_MEMBER]-> (p:Person), (forum) -[:CONTAINER_OF]-> (post:Post) -[:HAS_TAG]-> (tag:Tag {name: 'Manuel_Noriega'}) RETURN DISTINCT p.id AS PersonId LIMIT 200	MATCH (forum:Forum) -[:HAS_TAG]-> (tag:Tag {name: 'James_K._Polk'}) MATCH (forum) -[:HAS_MEMBER]-> (person:Person) WITH forum, COUNT(DISTINCT person) AS memberCount RETURN DISTINCT forum.title AS ForumTitle, memberCount ORDER BY memberCount DESC LIMIT 1	
	Structure: join_heavy, bounded_result.	Structure: join_heavy, aggregation, or- der_rank, bounded_result.	
	Relationships: CONTAINER_OF, HAS_MEMBER, HAS_TAG.	Relationships: HAS_MEMBER, HAS_TAG.	
	Gates: RO/Syn/Schema/Exec/Judge.	Gates: RO/Syn/Schema/Exec/Judge.	
	Result sample: {PersonId: 4398046511220}; observed rows: 19	Result sample: {ForumTitle: Wall of Javed Khan, memberCount: 33}; observed rows: 1	
12. SNB / Negation Difference / medium. Ques-	tion: Which members of forum 'Wall of Celso Oliveira' have not liked any post?	15. SNB / Simple Aggregation / easy. Question:	
	MATCH (forum:Forum {title: 'Wall of Celso Oliveira'}) -[:HAS_MEMBER]-> (p:Person) WHERE NOT (p) -[:LIKES]-> (:Post) RETURN DISTINCT p.id AS PersonId, p.firstName AS FirstName, p.lastName AS LastName LIMIT 200	MATCH (post:Post) -[:HAS_TAG]-> (tag:Tag {name: 'Vietnam'}) RETURN DISTINCT COUNT(DISTINCT post) AS PostCount	
	Structure: join_heavy, negation, bounded_result.	Structure: single_hop, aggregation.	
	Relationships: HAS_MEMBER, LIKES.	Relationships: HAS_TAG.	
	Gates: RO/Syn/Schema/Exec/Judge.	Gates: RO/Syn/Schema/Exec/Judge.	
	Result sample: {FirstName: K., LastName: Sen, PersonId: 94}; observed rows: 1	Result sample: {PostCount: 1}; observed rows: 1	
13. SNB / Path Temporal / easy. Question:	Which people are within two knows hops of person with id 4398046511136?	16. SNB / Simple Retrieval / easy. Ques- tion: Which post IDs did person with id 4398046511124 like?	
	MATCH (src:Person {id: 4398046511136}) -[:KNOWS*1..2]-> (dst:Person) RETURN DISTINCT dst.id AS PersonId, dst.firstName AS FirstName, dst.lastName	MATCH (p:Person {id: 4398046511124}) -[:LIKES]-> (post:Post) RETURN DISTINCT post.id AS PostId LIMIT 200	
	Structure: single_hop, bounded_result.	Structure: single_hop, bounded_result.	
	Relationships: LIKES.	Relationships: LIKES.	
	Gates: RO/Syn/Schema/Exec/Judge.	Gates: RO/Syn/Schema/Exec/Judge.	




Figure 2: Audited target-50 ablation suite over FinBench and SNB. The suite reaches all graph/variant/category targets and is included as appendix evidence while the larger target-100 and seeded repeat suites continue running for stronger final reliability evidence.




Figure 3: Gate-rate heatmap for the audited target-50 ablation suite. Deterministic read-only and syntax gates are saturated, while schema, execution, and judge rates expose the remaining quality differences across graph and pipeline variants.

Result sample: {PostId: 343597385744}; observed rows: 5




Figure 4: Diversity diagnostics for the full 3,000-example export. Category, graph-category, and difficulty balance are strong by construction; schema coverage and query-signature diversity expose remaining concentration from the seeded-template run.




Figure 5: Full 3,000-example export distribution. The benchmark preserves the planned 2:1 FinBench/SNB graph mix while balancing all eight Cypher workload categories and maintaining a near-even easy/medium difficulty split.




Figure 6: Downstream zero-shot Qwen3.5-9B Text2Cypher evaluation by workload category on the full test split. The breakdown separates final execution accuracy from parse/schema/execution success signals, making difficult operational categories visible.




Figure 7: Bootstrap uncertainty for downstream Qwen3.5-9B Text2Cypher evaluation by workload category. Error bars show 95% row-level bootstrap intervals over the full exported test split.




Figure 8: Downstream Text2Cypher failure taxonomy for the local Qwen3.5-9B baseline on the full test split. Exact matches are excluded so the chart exposes whether misses come from invalid Cypher, failed execution, or semantically wrong executable queries.